

# Data-Intensive Computing with MapReduce

Session 2: Hadoop Nuts and Bolts

Jimmy Lin  
University of Maryland  
Thursday, January 31, 2013



This work is licensed under a Creative Commons Attribution-Noncommercial-Share Alike 3.0 United States  
See <http://creativecommons.org/licenses/by-nc-sa/3.0/us/> for details





Source: Wikipedia (Japanese rock garden)



# **How will I actually learn Hadoop?**

- This class session
- Hadoop: The Definitive Guide
- RTFM
- RTFC(!)

# Materials in the course

- The course homepage
- Hadoop: The Definitive Guide
- Data-Intensive Text Processing with MapReduce
- Cloud<sup>9</sup>

**Take advantage of GitHub!**  
clone, branch, send pull request



# Basic Hadoop API\*

- Mapper

- `void setup(Mapper.Context context)`  
Called once at the beginning of the task
- `void map(K key, V value, Mapper.Context context)`  
Called once for each key/value pair in the input split
- `void cleanup(Mapper.Context context)`  
Called once at the end of the task

- Reducer/Combiner

- `void setup(Reducer.Context context)`  
Called once at the start of the task
- `void reduce(K key, Iterable<V> values, Reducer.Context context)`  
Called once for each key
- `void cleanup(Reducer.Context context)`  
Called once at the end of the task

\*Note that there are two versions of the API!

# Basic Hadoop API\*

- Partitioner

- int getPartition(K key, V value, int numPartitions)  
Get the partition number given total number of partitions

- Job

- Represents a packaged Hadoop job for submission to cluster
- Need to specify input and output paths
- Need to specify input and output formats
- Need to specify mapper, reducer, combiner, partitioner classes
- Need to specify intermediate/final key/value classes
- Need to specify number of reducers (but not mappers, why?)
- Don't depend of defaults!

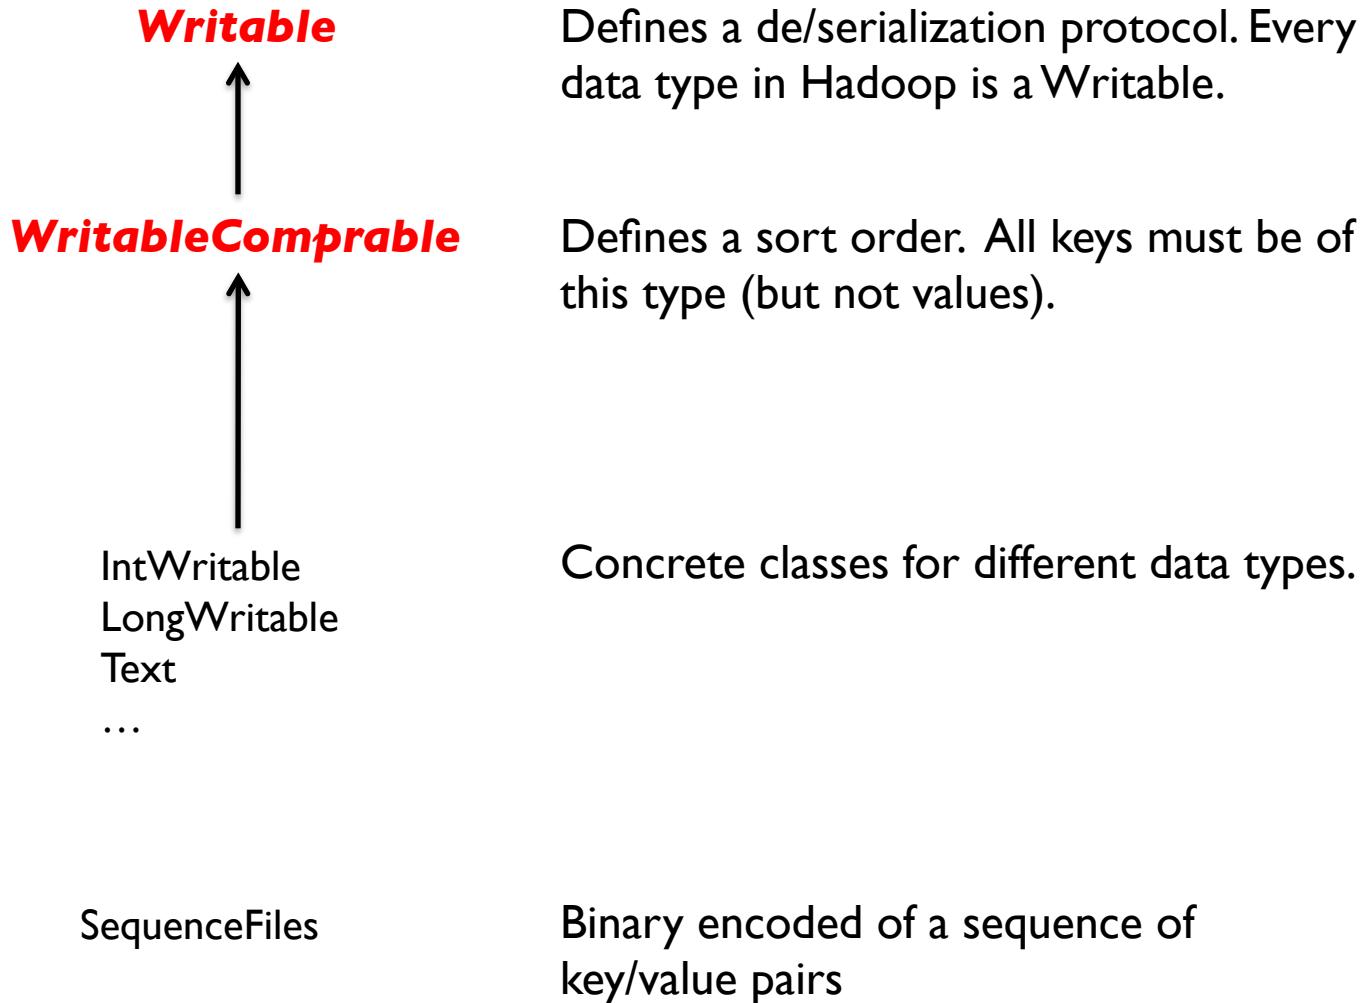
\*Note that there are two versions of the API!

# A tale of two packages...

org.apache.hadoop.mapreduce  
org.apache.hadoop.mapred



# Data Types in Hadoop: Keys and Values



# “Hello World”: Word Count

**Map(String docid, String text):**

for each word w in text:

    Emit(w, 1);

**Reduce(String term, Iterator<Int> values):**

    int sum = 0;

    for each v in values:

        sum += v;

    Emit(term, value);

# “Hello World”: Word Count

```
private static class MyMapper extends  
    Mapper<LongWritable, Text, Text, IntWritable> {  
  
    private final static IntWritable ONE = new IntWritable(1);  
    private final static Text WORD = new Text();  
  
    @Override  
    public void map(LongWritable key, Text value, Context context)  
        throws IOException, InterruptedException {  
        String line = ((Text) value).toString();  
        StringTokenizer itr = new StringTokenizer(line);  
        while (itr.hasMoreTokens()) {  
            WORD.set(itr.nextToken());  
            context.write(WORD, ONE);  
        }  
    }  
}
```

# “Hello World”: Word Count

```
private static class MyReducer extends  
    Reducer<Text, IntWritable, Text, IntWritable> {  
  
    private final static IntWritable SUM = new IntWritable();  
  
    @Override  
    public void reduce(Text key, Iterable<IntWritable> values,  
        Context context) throws IOException, InterruptedException {  
        Iterator<IntWritable> iter = values.iterator();  
        int sum = 0;  
        while (iter.hasNext()) {  
            sum += iter.next().get();  
        }  
        SUM.set(sum);  
        context.write(key, SUM);  
    }  
}
```

# Three Gotchas

- Avoid object creation at all costs
  - Reuse Writable objects, change the payload
- Execution framework reuses value object in reducer
- Passing parameters via class statics

# Getting Data to Mappers and Reducers

- Configuration parameters
  - Directly in the Job object for parameters
- “Side data”
  - DistributedCache
  - Mappers/reducers read from HDFS in setup method

# Complex Data Types in Hadoop

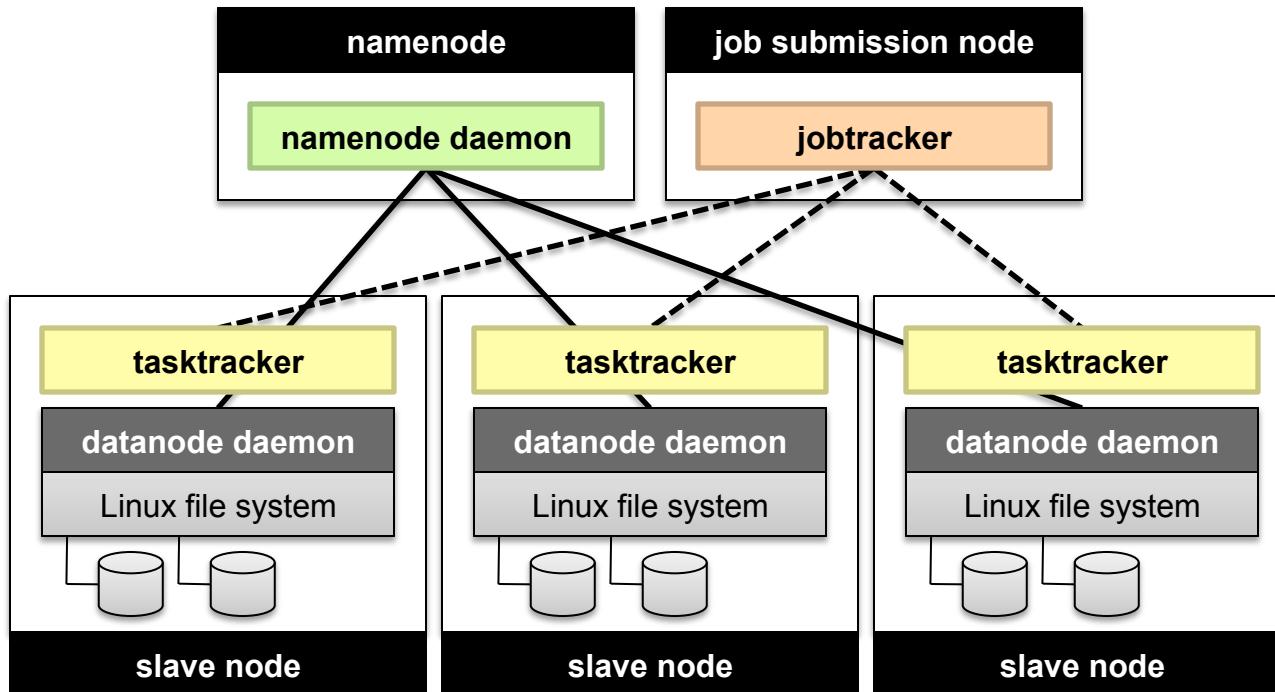
- How do you implement complex data types?
- The easiest way:
  - Encoded it as Text, e.g., (a, b) = “a:b”
  - Use regular expressions to parse and extract data
  - Works, but pretty hack-ish
- The hard way:
  - Define a custom implementation of Writable(Comparable)
  - Must implement: readFields, write, (compareTo)
  - Computationally efficient, but slow for rapid prototyping
  - Implement WritableComparator hook for performance
- Somewhere in the middle:
  - Cloud<sup>9</sup> offers JSON support and lots of useful Hadoop types
  - Quick tour...

# Basic Cluster Components\*

- One of each:
  - Namenode (NN): master node for HDFS
  - Jobtracker (JT): master node for job submission
- Set of each per slave machine:
  - Tasktracker (TT): contains multiple task slots
  - Datanode (DN): serves HDFS data blocks

\* Not quite... leaving aside YARN for now

# Putting everything together...



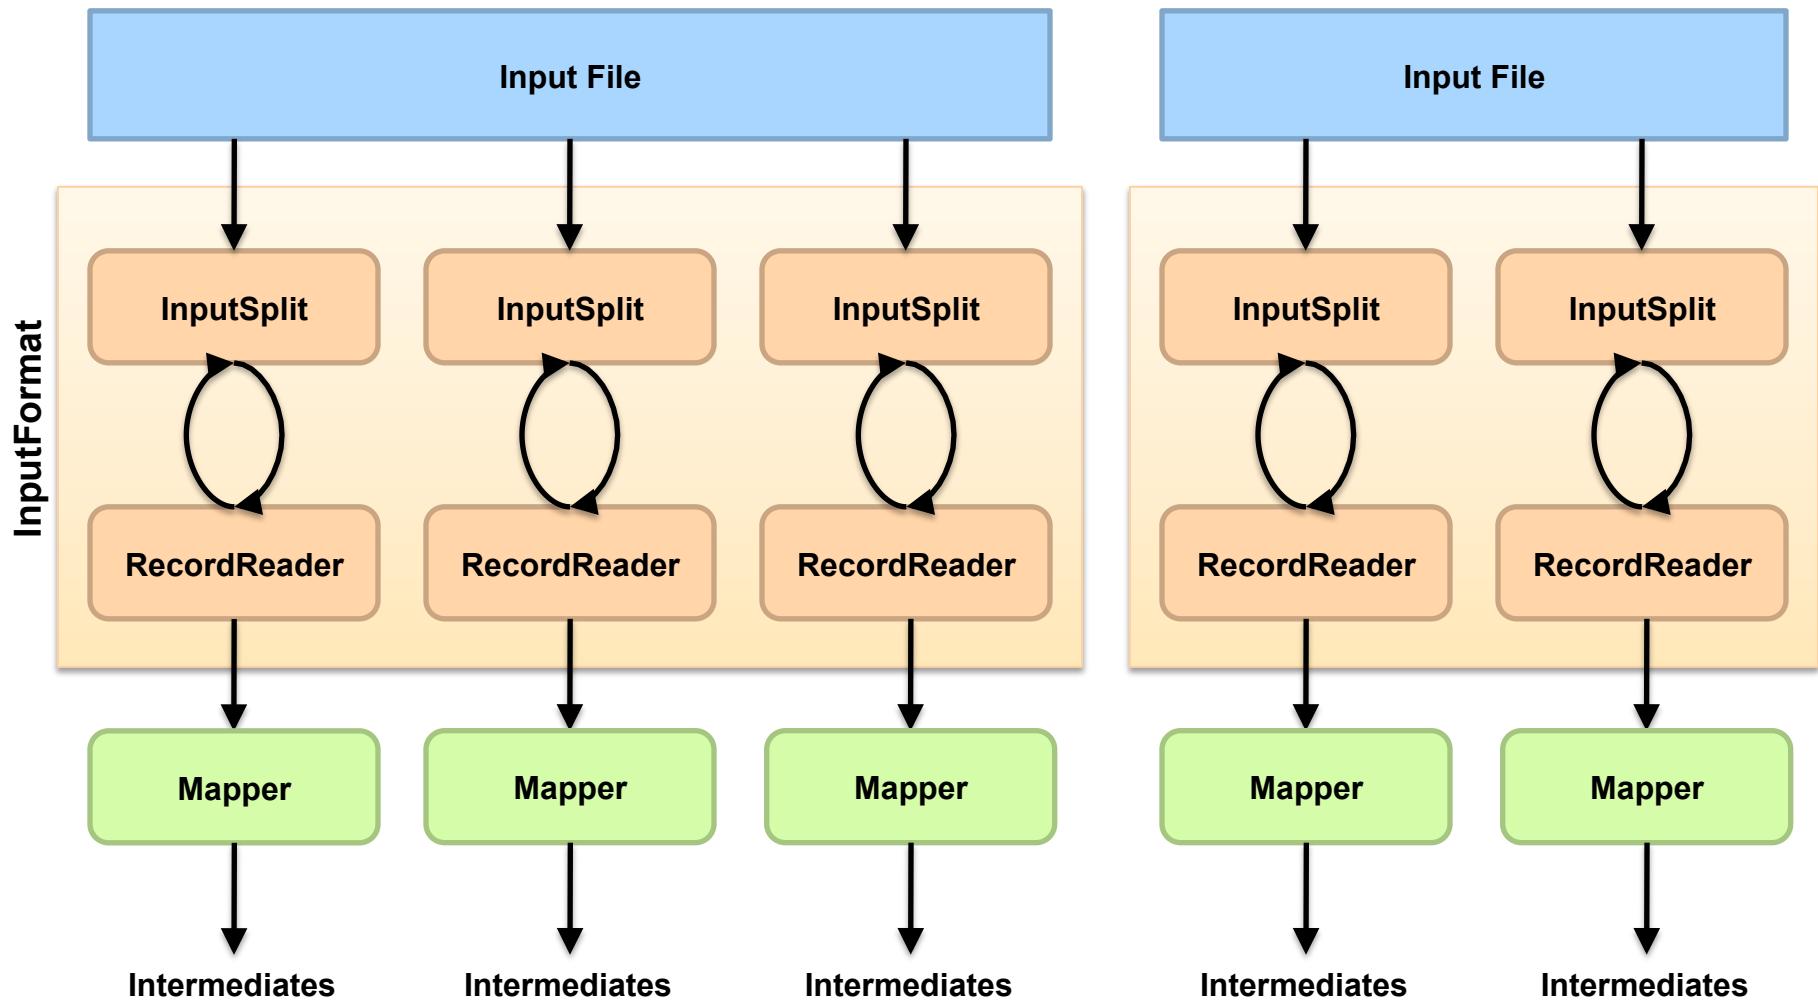
# Anatomy of a Job

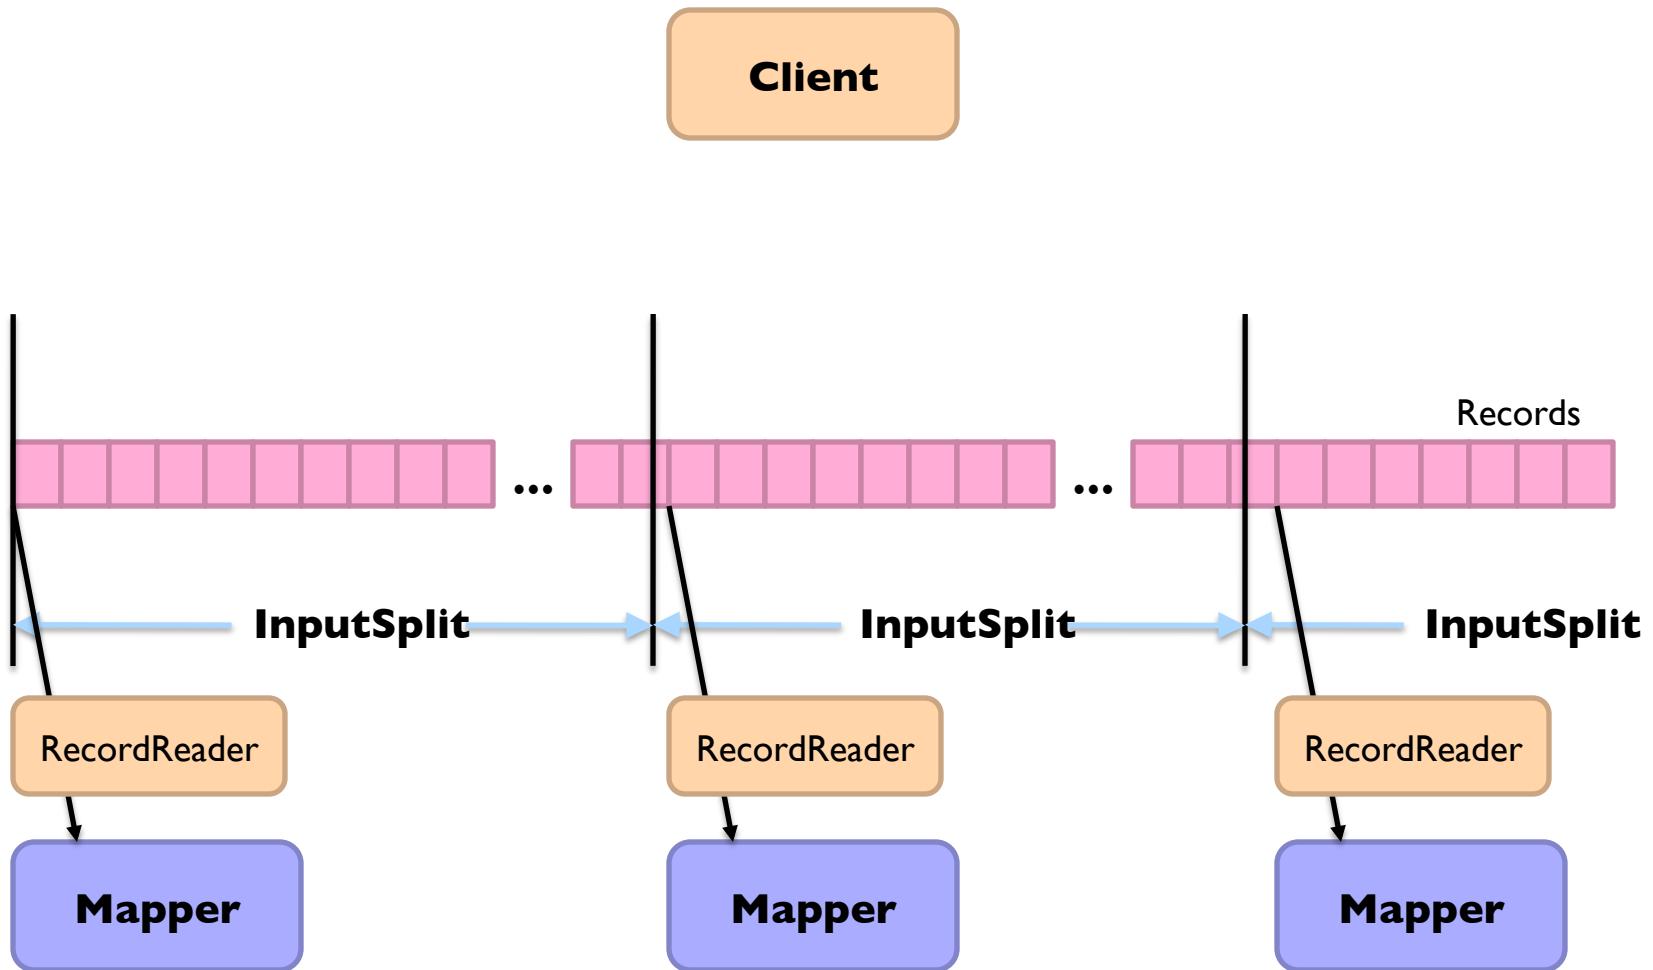
- MapReduce program in Hadoop = Hadoop job
  - Jobs are divided into map and reduce tasks
  - An instance of running a task is called a task attempt (occupies a slot)
  - Multiple jobs can be composed into a workflow
- Job submission:
  - Client (i.e., driver program) creates a job, configures it, and submits it to jobtracker
  - That's it! The Hadoop cluster takes over...

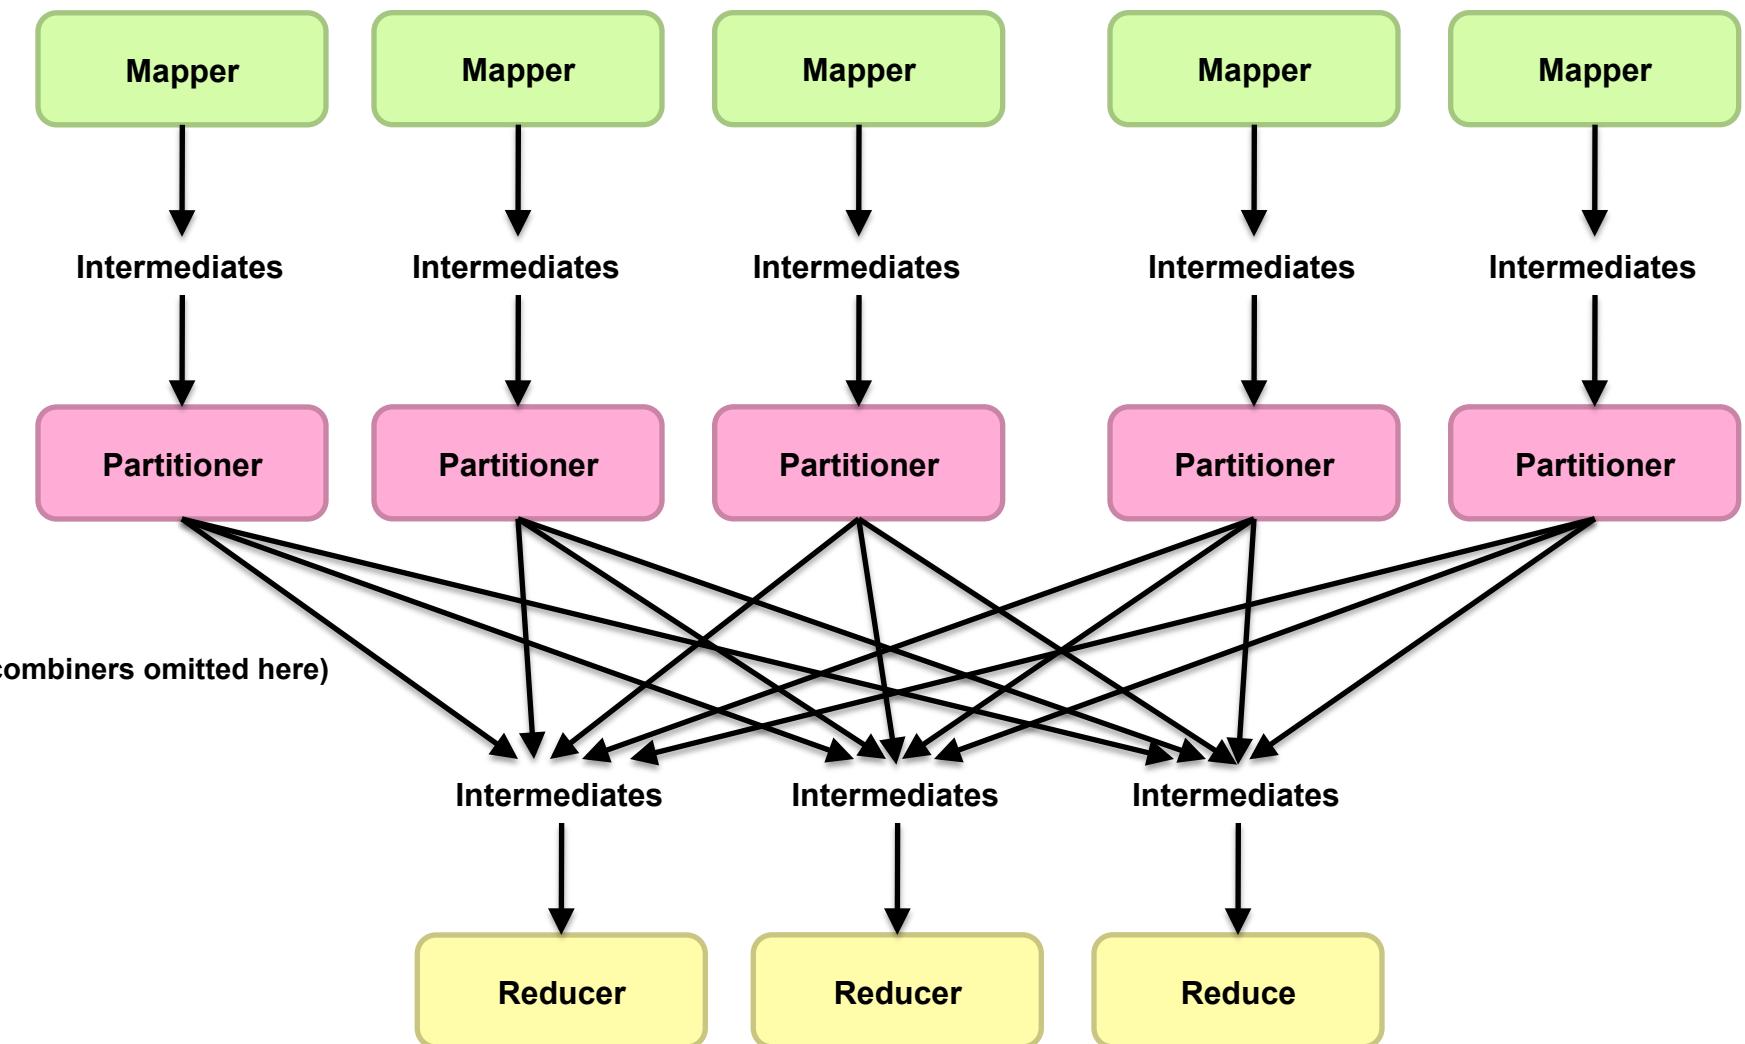
# Anatomy of a Job

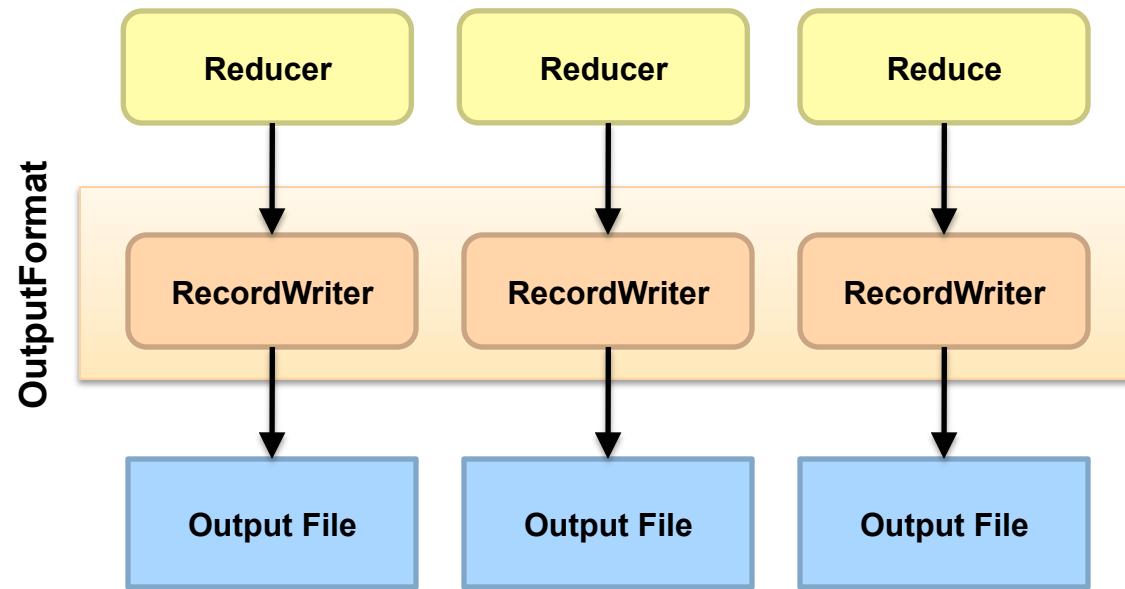
- Behind the scenes:

- Input splits are computed (on client end)
- Job data (jar, configuration XML) are sent to JobTracker
- JobTracker puts job data in shared location, enqueues tasks
- TaskTrackers poll for tasks
- Off to the races...









# **Input and Output**

- **InputFormat:**

- `TextInputFormat`
- `KeyValueTextInputFormat`
- `SequenceFileInputFormat`
- ...

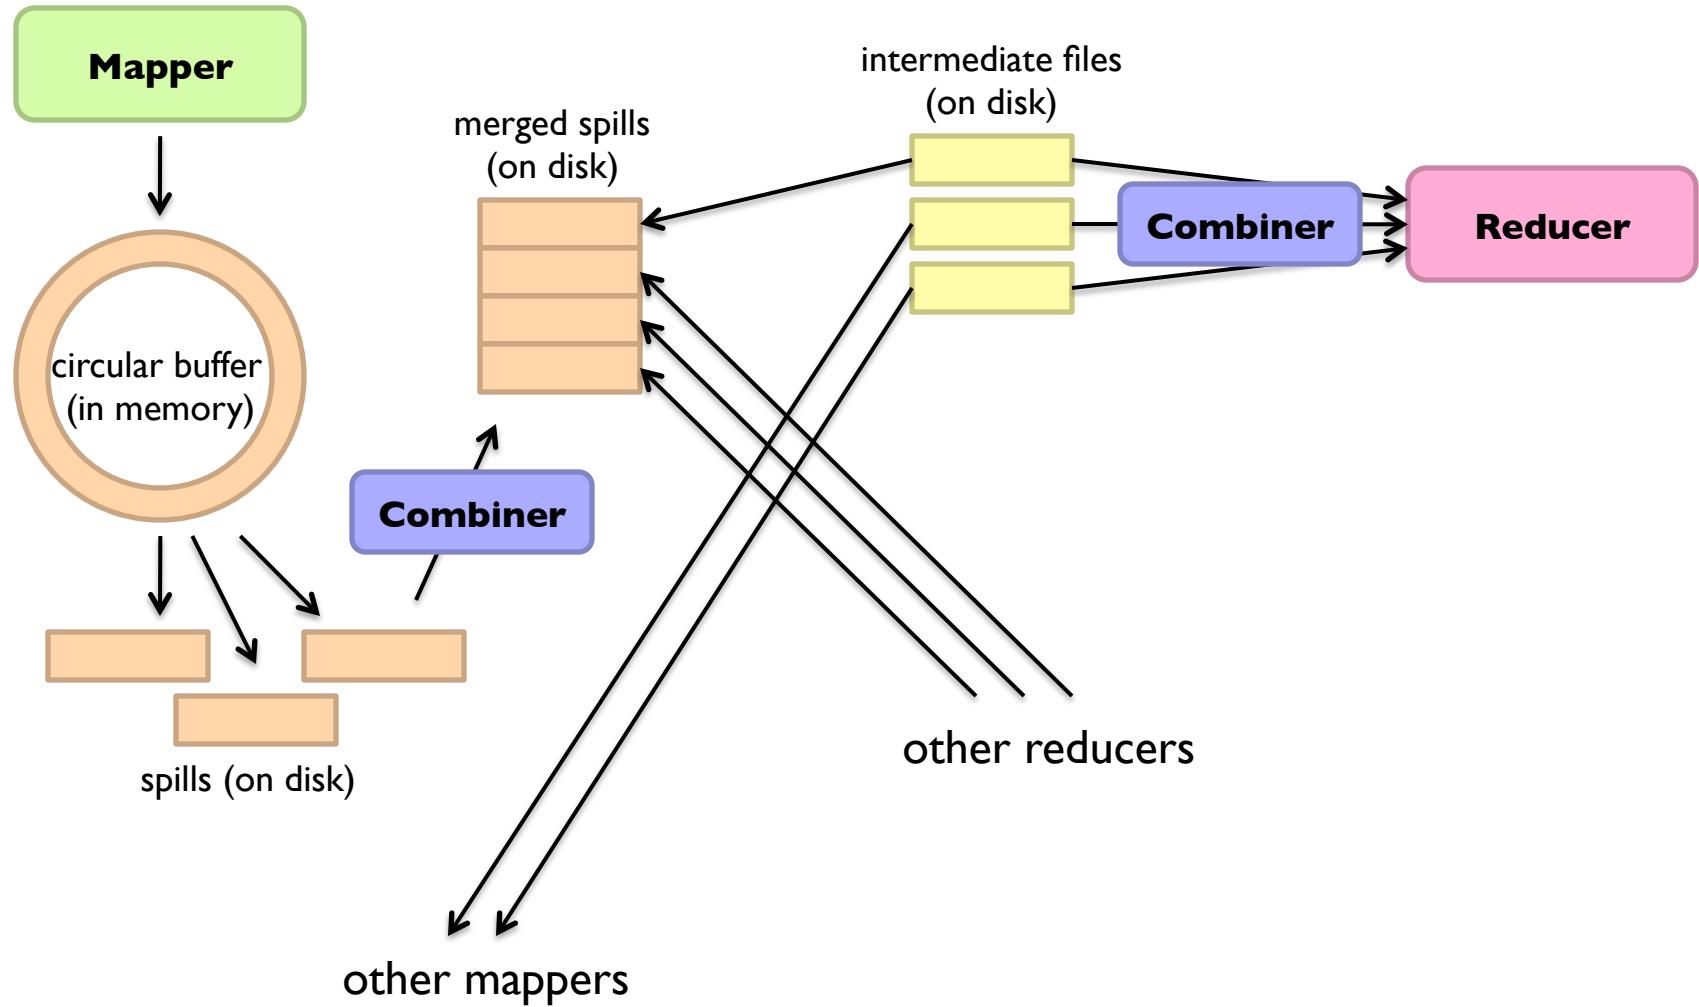
- **OutputFormat:**

- `TextOutputFormat`
- `SequenceFileOutputFormat`
- ...

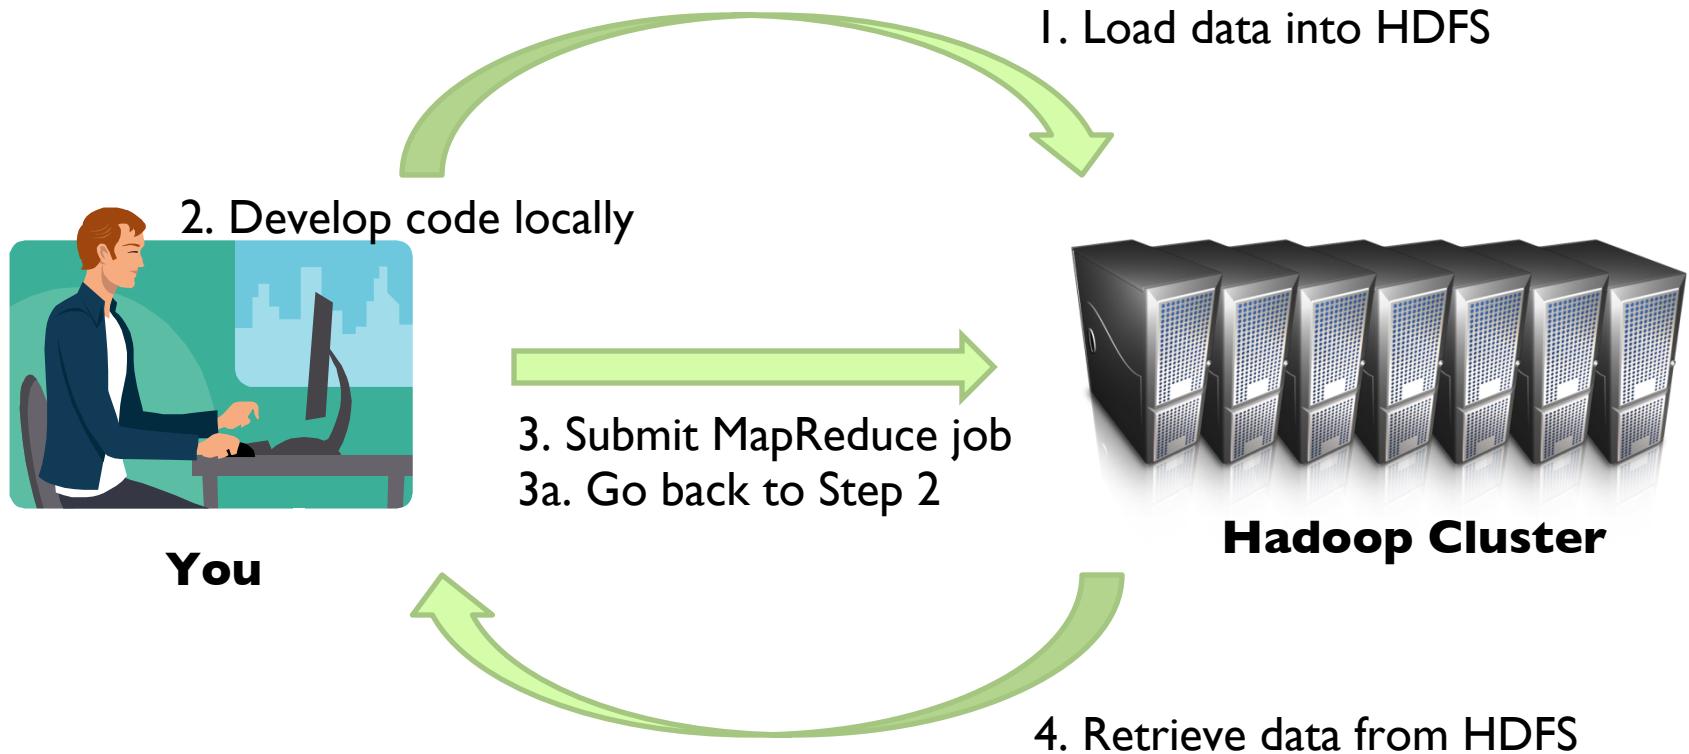
# Shuffle and Sort in Hadoop

- Probably the most complex aspect of MapReduce
- Map side
  - Map outputs are buffered in memory in a circular buffer
  - When buffer reaches threshold, contents are “spilled” to disk
  - Spills merged in a single, partitioned file (sorted within each partition): combiner runs during the merges
- Reduce side
  - First, map outputs are copied over to reducer machine
  - “Sort” is a multi-pass merge of map outputs (happens in memory and on disk): combiner runs during the merges
  - Final merge pass goes directly into reducer

# Shuffle and Sort



# Hadoop Workflow



# Recommended Workflow

- Here's how I work:
  - Develop code in Eclipse on host machine
  - Build distribution on host machine
  - Check out copy of code on VM
  - Copy (i.e., scp) jars over to VM (in same directory structure)
  - Run job on VM
  - Iterate
  - ...
  - Commit code on host machine and push
  - Pull from inside VM, verify
- Avoid using the UI of the VM
  - Directly ssh into the VM

# Actually Running the Job

- \$HADOOP\_CLASSPATH
- hadoop jar MYJAR.jar  
  -D k1=v1 ...  
  -libjars foo.jar,bar.jar  
  my.class.to.run arg1 arg2 arg3 ...

# Debugging Hadoop

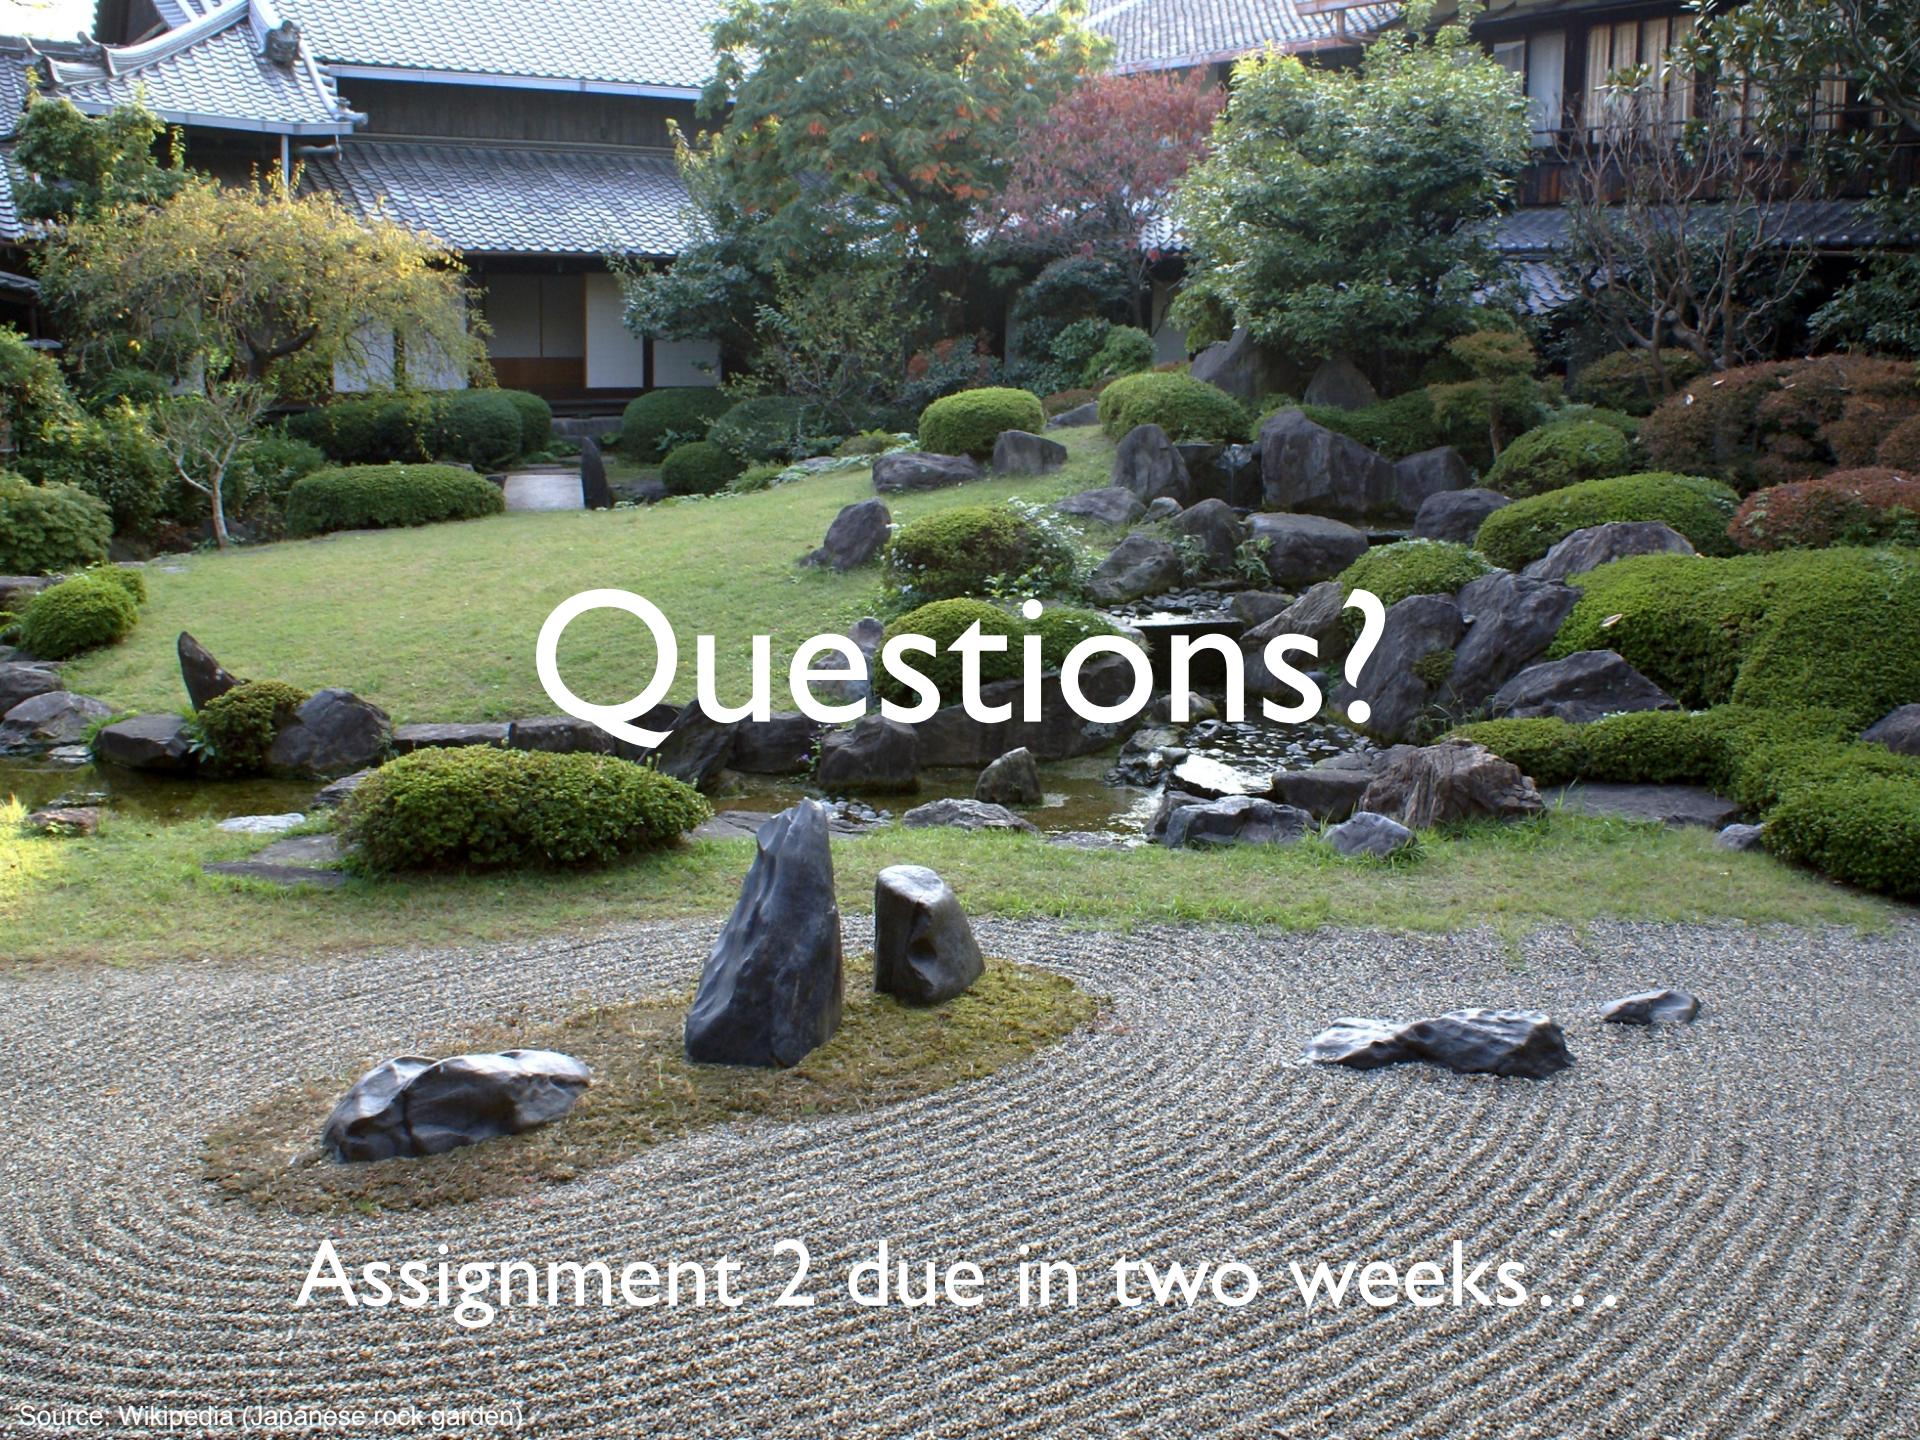
- First, take a deep breath
- Start small, start locally
- Build incrementally

# Code Execution Environments

- Different ways to run code:
  - Plain Java
  - Local (standalone) mode
  - Pseudo-distributed mode
  - Fully-distributed mode
- Learn what's good for what

# Hadoop Debugging Strategies

- Good ol' System.out.println
  - Learn to use the webapp to access logs
  - Logging preferred over System.out.println
  - Be careful how much you log!
- Fail on success
  - Throw RuntimeExceptions and capture state
- Programming is still programming
  - Use Hadoop as the “glue”
  - Implement core functionality outside mappers and reducers
  - Independently test (e.g., unit testing)
  - Compose (tested) components in mappers and reducers

A photograph of a traditional Japanese rock garden. The foreground features a gravel path with raked patterns. Several large, dark, irregular stones are scattered across the garden. In the middle ground, there is a small pond surrounded by rocks and low-lying green plants. The background consists of a building with a tiled roof and more manicured bushes and trees.

# Questions?

Assignment 2 due in two weeks...