

Data-Intensive Computing with MapReduce

Session 6: Similar Item Detection

Jimmy Lin
University of Maryland
Thursday, February 28, 2013



This work is licensed under a Creative Commons Attribution-Noncommercial-Share Alike 3.0 United States
See <http://creativecommons.org/licenses/by-nc-sa/3.0/us/> for details

Developing “big data” intuitions...



Today's Agenda

- The problem of similar item detection
- Distances and representations
- Minhash
- Random projections
- End-to-end Wikipedia application

What's the Problem?

- Finding similar items with respect to some distance metric
- Two variants of the problem:
 - Offline: extract all similar pairs of objects from a large collection
 - Online: is this object similar to something I've seen before?
- Similarities/differences with the clustering problem

Applications

- Near-duplicate detection of webpages
 - Offline vs. online variant of problem
- Recommender systems
- Forensic text analysis (e.g., plagiarism detection)

Near-Duplicate Detection of Webpages

- What's the source of the problem?
 - Mirror pages (legit)
 - Spam farms (non-legit)
 - Additional complications (e.g., nav bars)
- Naïve algorithm:
 - Compute cryptographic hash for webpage (e.g., MD5)
 - Insert hash values into a big hash table
 - Compute hash for new webpage: collision implies duplicate
- What's the issue?
- Intuition:
 - Hash function needs to be tolerant of minor differences
 - High similarity implies higher probability of hash collision

Literature Note

- Many communities have tackled similar problems:
 - Theoretical computer science
 - Information retrieval
 - Data mining
 - Databases
 - ...
- Issues
 - Slightly different terminology
 - Results not easy to compare

Four Steps

- Specify distance metric
 - Jaccard, Euclidean, cosine, etc.
- Compute representation
 - Shingling, tf.idf, etc.
- “Project”
 - Minhash, random projections, etc.
- Extract
 - Bucketing, sliding windows, etc.

Distances



Distance Metrics

1. Non-negativity:

$$d(x, y) \geq 0$$

2. Identity:

$$d(x, y) = 0 \iff x = y$$

3. Symmetry:

$$d(x, y) = d(y, x)$$

4. Triangle Inequality

$$d(x, y) \leq d(x, z) + d(z, y)$$

Distance: Jaccard

- Given two sets A, B
- Jaccard similarity:

$$J(A, B) = \frac{|A \cap B|}{|A \cup B|}$$

$$d(A, B) = 1 - J(A, B)$$

Distance: Norms

- Given: $x = [x_1, x_2, \dots, x_n]$
 $y = [y_1, y_2, \dots, y_n]$

- Euclidean distance (L_2 -norm)

$$d(x, y) = \sqrt{\sum_{i=0}^n (x_i - y_i)^2}$$

- Manhattan distance (L_1 -norm)

$$d(x, y) = \sum_{i=0}^n |x_i - y_i|$$

- L_r -norm

$$d(x, y) = \left[\sum_{i=0}^n |x_i - y_i|^r \right]^{1/r}$$

Distance: Cosine

- Given: $x = [x_1, x_2, \dots, x_n]$
 $y = [y_1, y_2, \dots, y_n]$
- Idea: measure distance between the vectors

$$\cos \theta = \frac{\mathbf{x} \cdot \mathbf{y}}{|\mathbf{x}| |\mathbf{y}|}$$

- Thus:

$$\text{sim}(\mathbf{x}, \mathbf{y}) = \frac{\sum_{i=0}^n x_i y_i}{\sqrt{\sum_{i=0}^n x_i^2} \sqrt{\sum_{i=0}^n y_i^2}}$$

$$d(\mathbf{x}, \mathbf{y}) = 1 - \text{sim}(\mathbf{x}, \mathbf{y})$$

Distance: Hamming

- Given two bit vectors
- Hamming distance: number of elements which differ

Representations



Representations: Text

- Unigrams (i.e., words)
- Shingles = n -grams
 - At the word level
 - At the character level
- Feature weights
 - boolean
 - tf.idf
 - BM25
 - ...
- Which representation to use?
- What's the actual text?

Use the hashing trick!

Representations: Beyond Text

- For recommender systems:
 - Items as features for users
 - Users as features for items
- For graphs:
 - Adjacency lists as features for vertices
- With log data:
 - Behaviors (clicks) as features

Minhash



Minhash

- Seminal algorithm for near-duplicate detection of webpages
 - Used by AltaVista
 - For details see Broder et al. (1997)
- Setup:
 - Documents (HTML pages) represented by shingles (n -grams)
 - Jaccard similarity: dups are pairs with high similarity

Preliminaries: Representation

- Sets:

- $A = \{e_1, e_3, e_7\}$
- $B = \{e_3, e_5, e_7\}$

- Can be equivalently expressed as matrices:

Element	A	B
e_1	1	0
e_2	0	0
e_3	1	1
e_4	0	0
e_5	0	1
e_6	0	0
e_7	1	1

Preliminaries: Jaccard

Element	A	B
e ₁	1	0
e ₂	0	0
e ₃	1	1
e ₄	0	0
e ₅	0	1
e ₆	0	0
e ₇	1	1

Let:

M_{00} = # rows where both elements are 0

M_{11} = # rows where both elements are 1

M_{01} = # rows where A=0, B=1

M_{10} = # rows where A=1, B=0

$$J(A, B) = \frac{M_{11}}{M_{01} + M_{10} + M_{11}}$$

Minhash

- Computing minhash

- Start with the matrix representation of the set
- Randomly permute the rows of the matrix
- minhash is the first row with a “one”

- Example:

$$h(A) = e_3 \quad h(B) = e_5$$

Element	A	B
e_1	1	0
e_2	0	0
e_3	1	1
e_4	0	0
e_5	0	1
e_6	0	0
e_7	1	1

Element	A	B
e_6	0	0
e_2	0	0
e_5	0	1
e_3	1	1
e_7	1	1
e_4	0	0
e_1	1	0

Minhash and Jaccard

Element	A	B	
e ₆	0	0	M ₀₀
e ₂	0	0	M ₀₀
e ₅	0		M ₀₁
e ₃			M ₁₁
e ₇			M ₁₁
e ₄	0	0	M ₀₀
e ₁		0	M ₁₀

$$P[h(A) = h(B)] = \text{J}(A, B)$$

$$\frac{M_{11}}{M_{01} + M_{10} + M_{11}} \quad \frac{M_{11}}{M_{01} + M_{10} + M_{11}}$$

To Permute or Not to Permute?

- Permutations are expensive
- Interpret the hash value as the permutation
- Only need to keep track of the minimum hash value
 - Can keep track of multiple minhash values at once

Extracting Similar Pairs (LSH)

- We know: $P[h(A) = h(B)] = J(A, B)$
- Task: discover all pairs with similarity greater than s
- Algorithm:
 - For each object, compute its minhash value
 - Group objects by their hash values
 - Output all pairs within each group
- Analysis:
 - Probability we will discovered all pairs is s
 - Probability that any pair is invalid is $(1 - s)$
- What's the fundamental issue?

Two Minhash Signatures

- Task: discover all pairs with similarity greater than s
- Algorithm:
 - For each object, compute two minhash values and concatenate together into a signature
 - Group objects by their signatures
 - Output all pairs within each group
- Analysis:
 - Probability we will discovered all pairs is s^2
 - Probability that any pair is invalid is $(1 - s)^2$

k Minhash Signatures

- Task: discover all pairs with similarity greater than s
- Algorithm:
 - For each object, compute k minhash values and concatenate together into a signature
 - Group objects by their signatures
 - Output all pairs within each group
- Analysis:
 - Probability we will discovered all pairs is s^k
 - Probability that any pair is invalid is $(1 - s)^k$
- What's the issue now?

n different k Minhash Signatures

- Task: discover all pairs with similarity greater than s
- Algorithm:
 - For each object, compute n sets k minhash values
 - For each set, concatenate k minhash values together
 - Within each set:
 - Group objects by their signatures
 - Output all pairs within each group
 - De-dup pairs
- Analysis:
 - Probability we will miss a pair is $(1 - s^k)^n$
 - Probability that any pair is invalid is $n(1 - s)^k$

Practical Notes

- In some cases, checking all candidate pairs may be possible
 - Time cost is small relative to everything else
 - Easy method to discard false positives
- Most common practical implementation:
 - Generate M minhash values, randomly select k of them n times
 - Reduces amount of hash computations needed
- Determining “authoritative” version is non-trivial

MapReduce Implementation

- Map over objects:
 - Generate M minhash values, randomly select k of them n times
 - Each draw yields a signature: emit as intermediate key, value is object id
- Shuffle/sort:
- Reduce:
 - Receive all object ids with same signature, generate candidate pairs
 - Okay to buffer in memory?
- Second pass to de-dup

Offline Extraction vs. Online Querying

- Batch formulation of the problem:
 - Discover all pairs with similarity greater than s
 - Useful for post-hoc batch processing of web crawl
- Online formulation of the problem:
 - Given new webpage, is it similar to one I've seen before?
 - Useful for incremental web crawl processing

Online Similarity Querying

- Preparing the existing collection:
 - For each object, compute n sets of k minhash values
 - For each set, concatenate k minhash values together
 - Keep each signature in hash table (in memory)
 - Note: can parallelize across multiple machines
- Querying and updating:
 - For new webpage, compute signatures and check for collisions
 - Collisions imply duplicate (determine which version to keep)
 - Update hash tables

Random Projections

Limitations of Minhash

- Minhash is great for near-duplicate detection
 - Set high threshold for Jaccard similarity
- Limitations:
 - Jaccard similarity only
 - Set-based representation, no way to assign weights to features
- Random projections:
 - Works with arbitrary vectors using cosine similarity
 - Same basic idea, but details differ
 - Slower but more accurate: no free lunch!

Random Projection Hashing

- Generate a random vector r of unit length
 - Draw from univariate Gaussian for each component
 - Normalize length
- Define:

$$h_r(u) = \begin{cases} 1 & \text{if } r \cdot u \geq 0 \\ 0 & \text{if } r \cdot u < 0 \end{cases}$$

- Physical intuition?

RP Hash Collisions

- It can be shown that:

$$P[h_r(u) = h_r(v)] = 1 - \frac{\theta(u, v)}{\pi}$$

- Proof in (Goemans and Williamson, 1995)
- Physical intuition?

Random Projection Signature

- Given D random vectors:

$$[r_1, r_2, r_3, \dots, r_D]$$

- Convert each object into a D bit signature

$$u \rightarrow [h_{r_1}(u), h_{r_2}(u), h_{r_3}(u), \dots, h_{r_D}(u)]$$

- We can derive:

$$\text{sim}(u, v) = \cos \left[\left(\frac{\pi}{D} \right) \text{hamming}(s_u, s_v) \right]$$

- Thus: similarity boils down to comparison of hamming distances between signatures

One-RP Signature

- Task: discover all pairs with cosine similarity greater than s
- Algorithm:
 - Compute D -bit RP signature for every object
 - Take first bit, bucket objects into two sets
 - Perform brute force pairwise (hamming distance) comparison in each bucket, retain those below hamming distance threshold
- Analysis:
 - Probability we will discover all pairs:^{*}
$$1 - \frac{\cos^{-1}(s)}{\pi}$$
 - Efficiency:

$$N^2 \quad \text{vs.} \quad 2 \left(\frac{N}{2} \right)^2$$

* Note, this is actually a simplification: see Ture et al. (SIGIR 2011) for details.

Two-RP Signature

- Task: discover all pairs with cosine similarity greater than s
- Algorithm:
 - Compute D -bit RP signature for every object
 - Take first two bits, bucket objects into four sets
 - Perform brute force pairwise (hamming distance) comparison in each bucket, retain those below hamming distance threshold
- Analysis:
 - Probability we will discover all pairs:
$$\left[1 - \frac{\cos^{-1}(s)}{\pi}\right]^2$$
 - Efficiency:

$$N^2 \quad \text{vs.} \quad 4 \left(\frac{N}{4}\right)^2$$

k-RP Signature

- Task: discover all pairs with cosine similarity greater than s
- Algorithm:
 - Compute D -bit RP signature for every object
 - Take first k bits, bucket objects into 2^k sets
 - Perform brute force pairwise (hamming distance) comparison in each bucket, retain those below hamming distance threshold
- Analysis:
 - Probability we will discover all pairs:
$$\left[1 - \frac{\cos^{-1}(s)}{\pi}\right]^k$$
 - Efficiency:
$$N^2 \quad \text{vs.} \quad 2^k \left(\frac{N}{2^k}\right)^2$$

m Sets of k -RP Signature

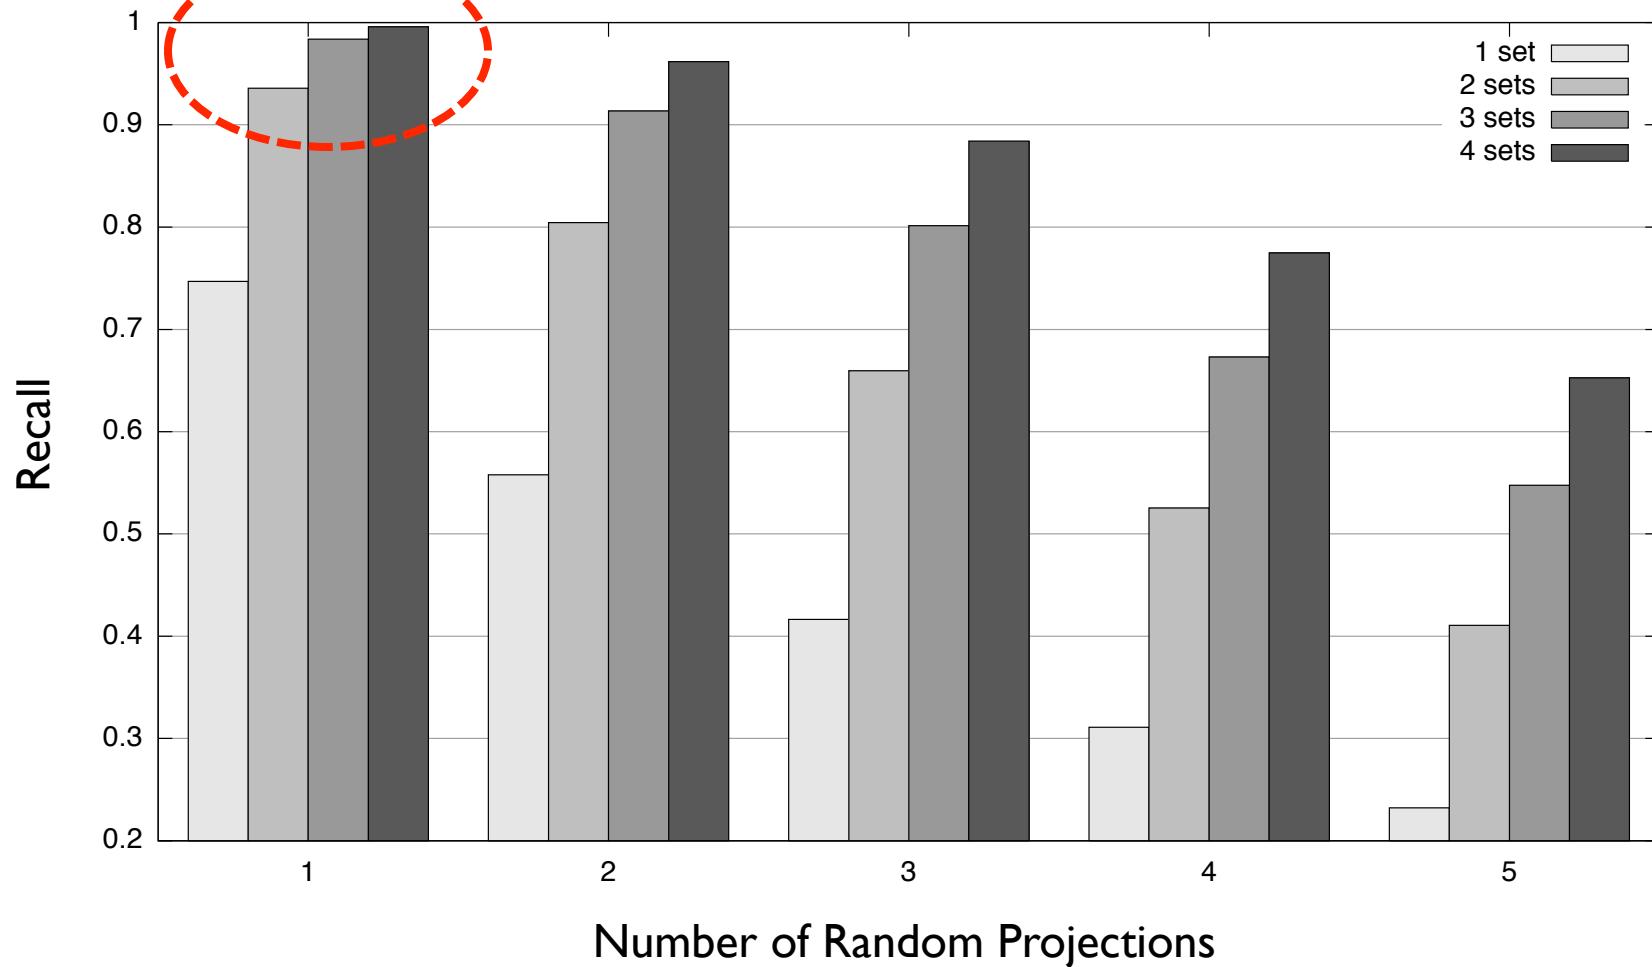
- Task: discover all pairs with cosine similarity greater than s
- Algorithm:
 - Compute D -bit RP signature for every object
 - Choose m sets of k bits
 - For each set, use k selected bits to partition objects into 2^k sets
 - Perform brute force pairwise (hamming distance) comparison in each bucket (of each set), retain those below hamming distance threshold
- Analysis:
 - Probability we will discover all pairs:
$$1 - \left[1 - \left[1 - \frac{\cos^{-1}(s)}{\pi} \right]^k \right]^m$$
 - Efficiency: N^2 vs. $m \cdot 2^k \left(\frac{N}{2^k} \right)^2$

Theoretical Results

Threshold = 0.7

What's interesting here?

Recall



Online Querying

- Preprocessing:
 - Compute D -bit RP signature for every object
 - Choose m sets of k bits and use to bucket
 - Store signatures in memory (across multiple machines)
- Querying
 - Compute D -bit signature of query object, choose m sets of k bits in same way
 - Perform brute-force scan of correct bucket (in parallel)

Additional Issues to Consider

- Emphasis on recall, not precision
- Two sources of error:
 - From LSH
 - From using hamming distance as proxy for cosine similarity
- Balance of bucket sizes?
- Parameter tuning?

“Sliding Window” Algorithm

- Compute D -bit RP signature for every object
- For each object, permute bit signature m times
- For each permutation, sort bit signatures
 - Apply sliding window of width B over sorted
 - Compute hamming distances of bit signatures within window

MapReduce Implementation

- Mapper:

- Process each individual object in parallel
- Load in random vectors as side data
- Compute bit signature
- Permute m times, for each emit:
key = $(n, \text{signature})$, where $n = [1 \dots m]$
value = object id

- Reduce

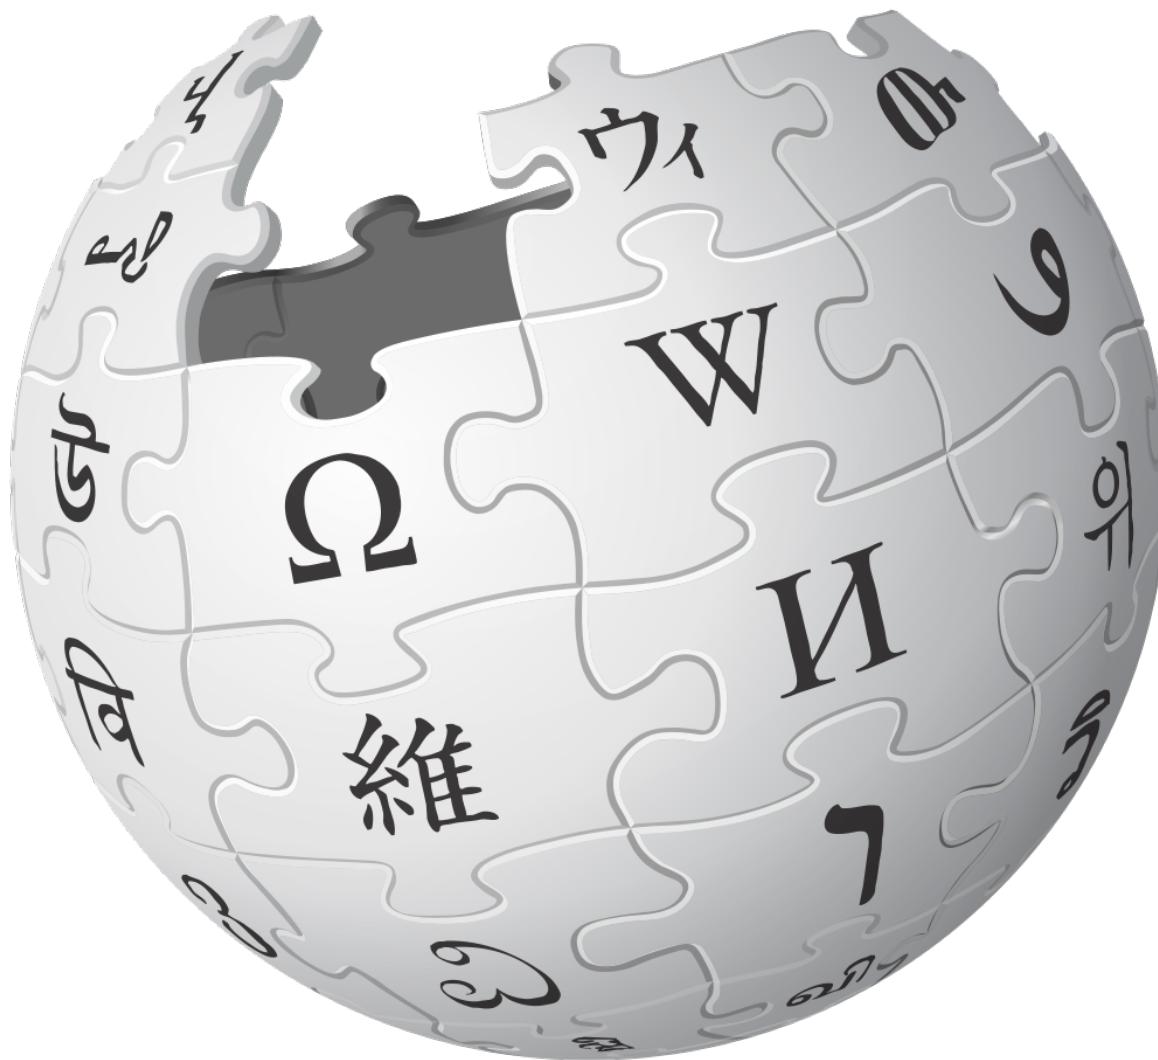
- Keep FIFO queue of B bit signatures
- For each newly-encountered bit signature, compute hamming distance wrt all bit signatures in queue
- Add new bit signature to end of queue, displacing oldest

Slightly Faster MapReduce Variant

- In reducer, don't perform hamming distance computations
 - Write directly to HDFS in large blocks
- Perform hamming distance computations in second mapper-only pass
 - Advantages?

Case Study: Wikipedia

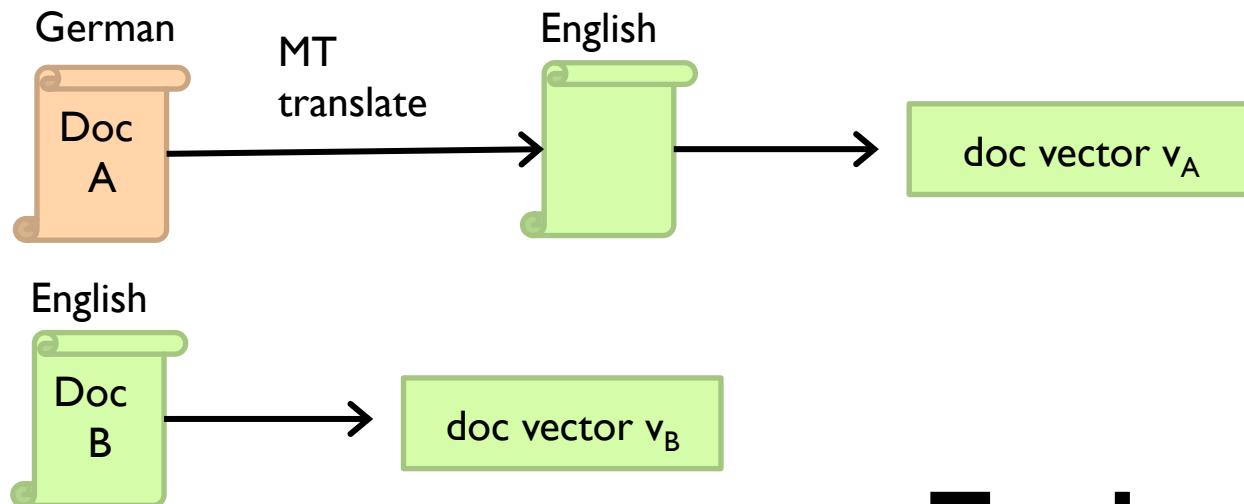
Ture et al., SIGIR 2011
Ture and Lin, NAACL 2012



Mining Bitext from Wikipedia

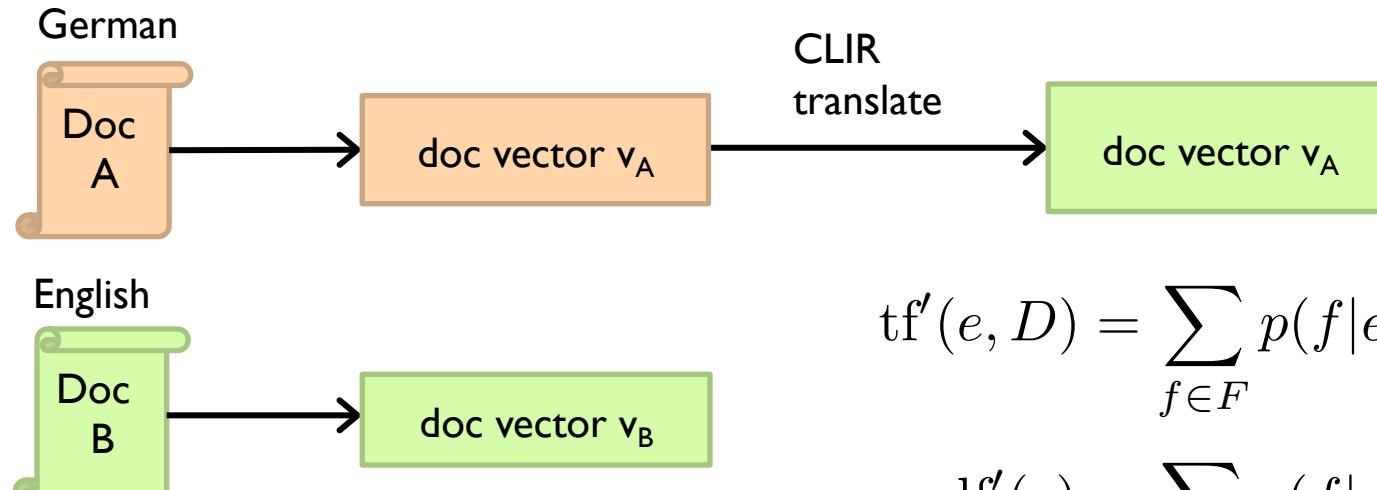
- Task: extract similar article pairs in *different* languages
 - Why is this useful?
 - Why not just use inter-wiki links?
- Use LSH!
 - Why is this non-trivial?

Machine Translation



Tradeoffs?

“Vector” Translation

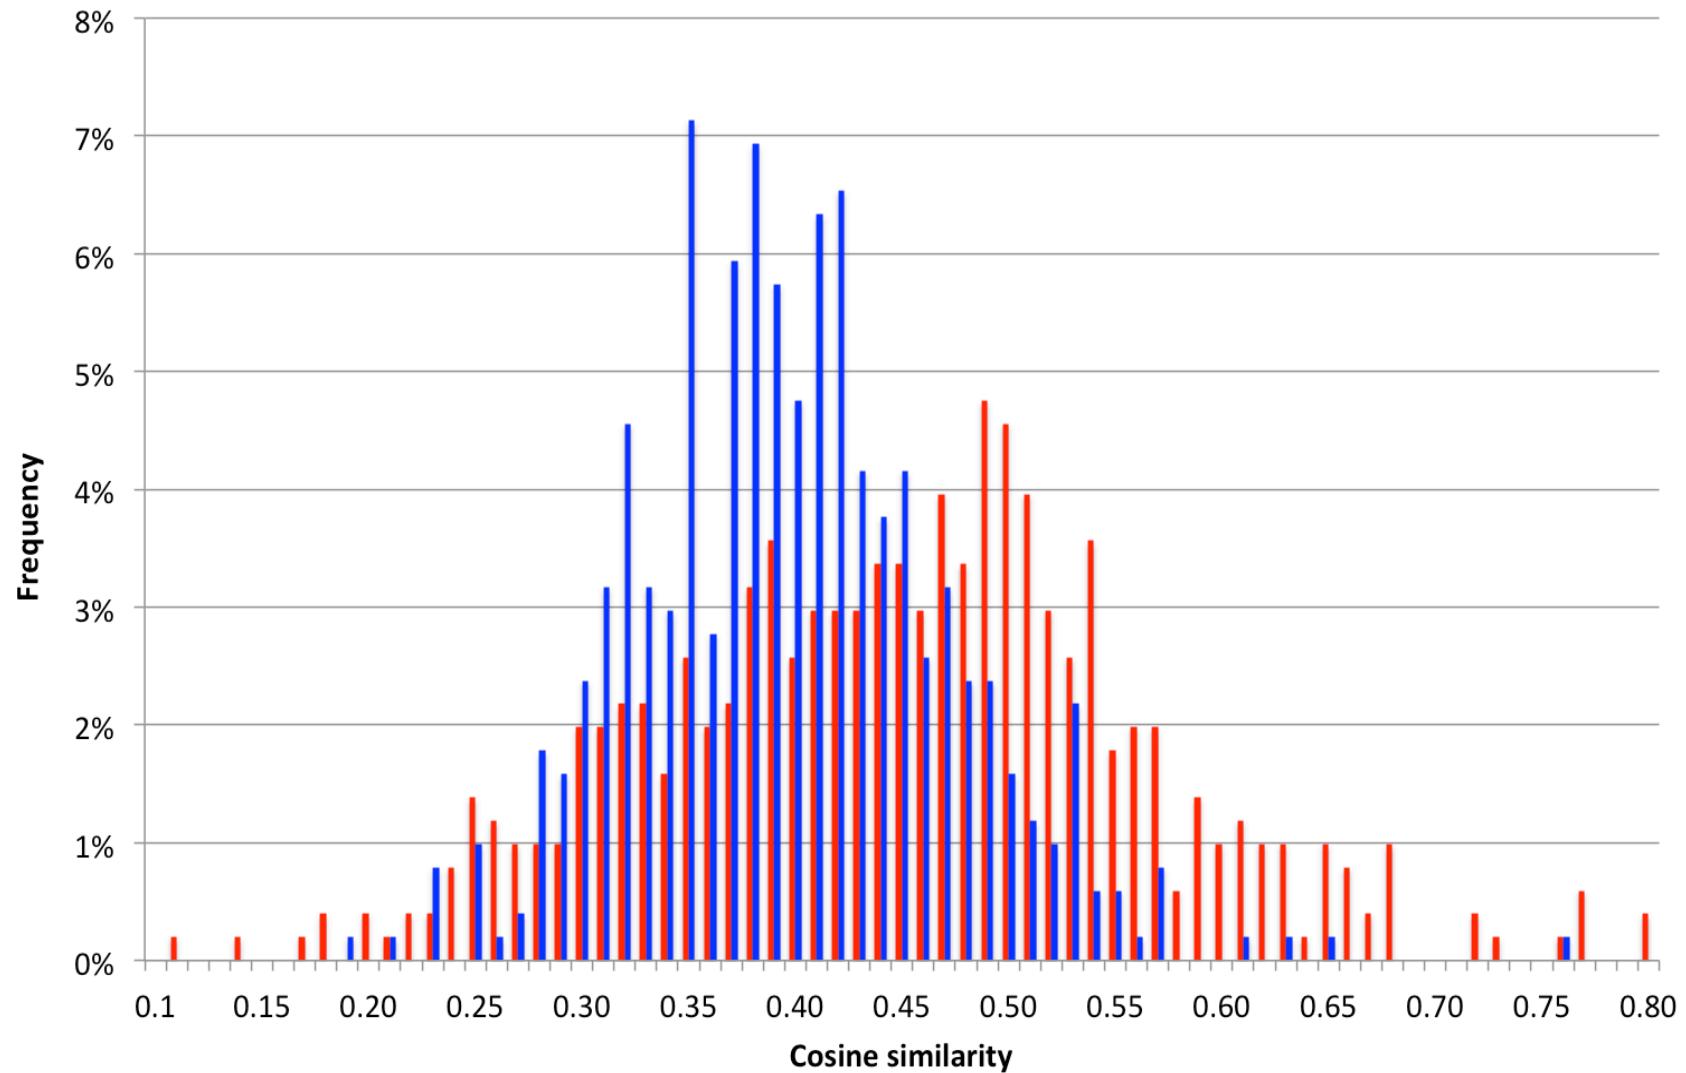


$$tf'(e, D) = \sum_{f \in F} p(f|e) \cdot tf(f, D)$$

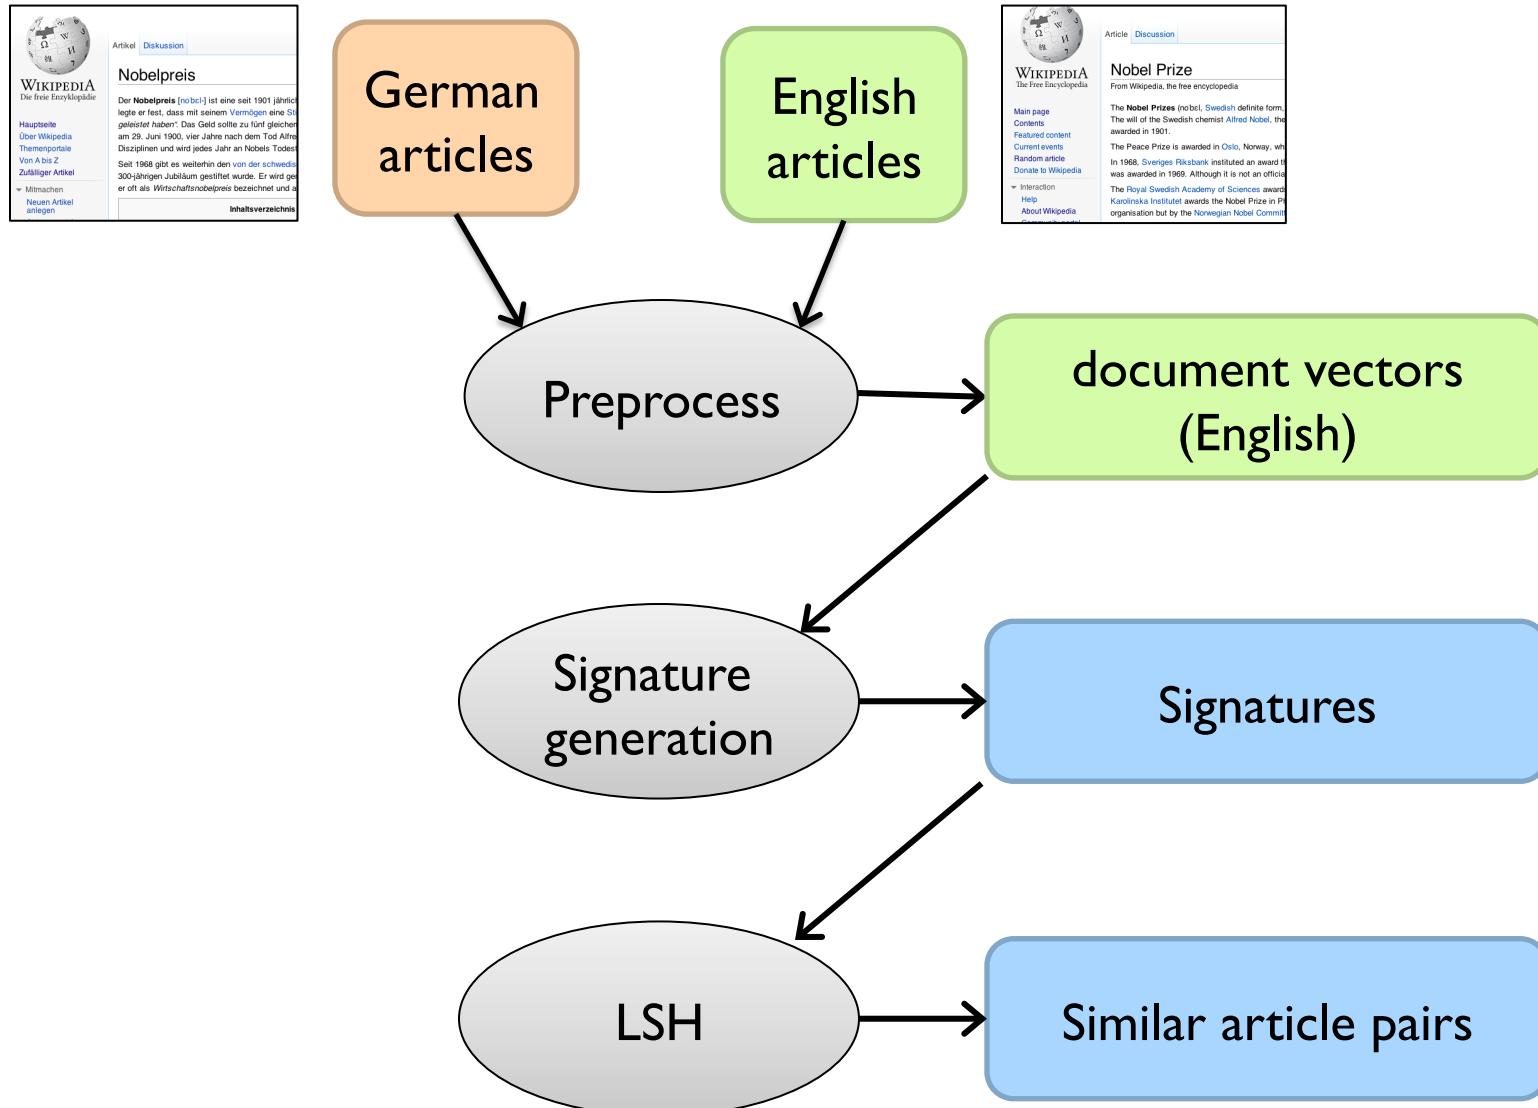
$$df'(e) = \sum_{f \in F} p(f|e) \cdot df(f)$$

Translations are noisy!

Question: To what extent are “ground truth” mutual translation pairs similar?



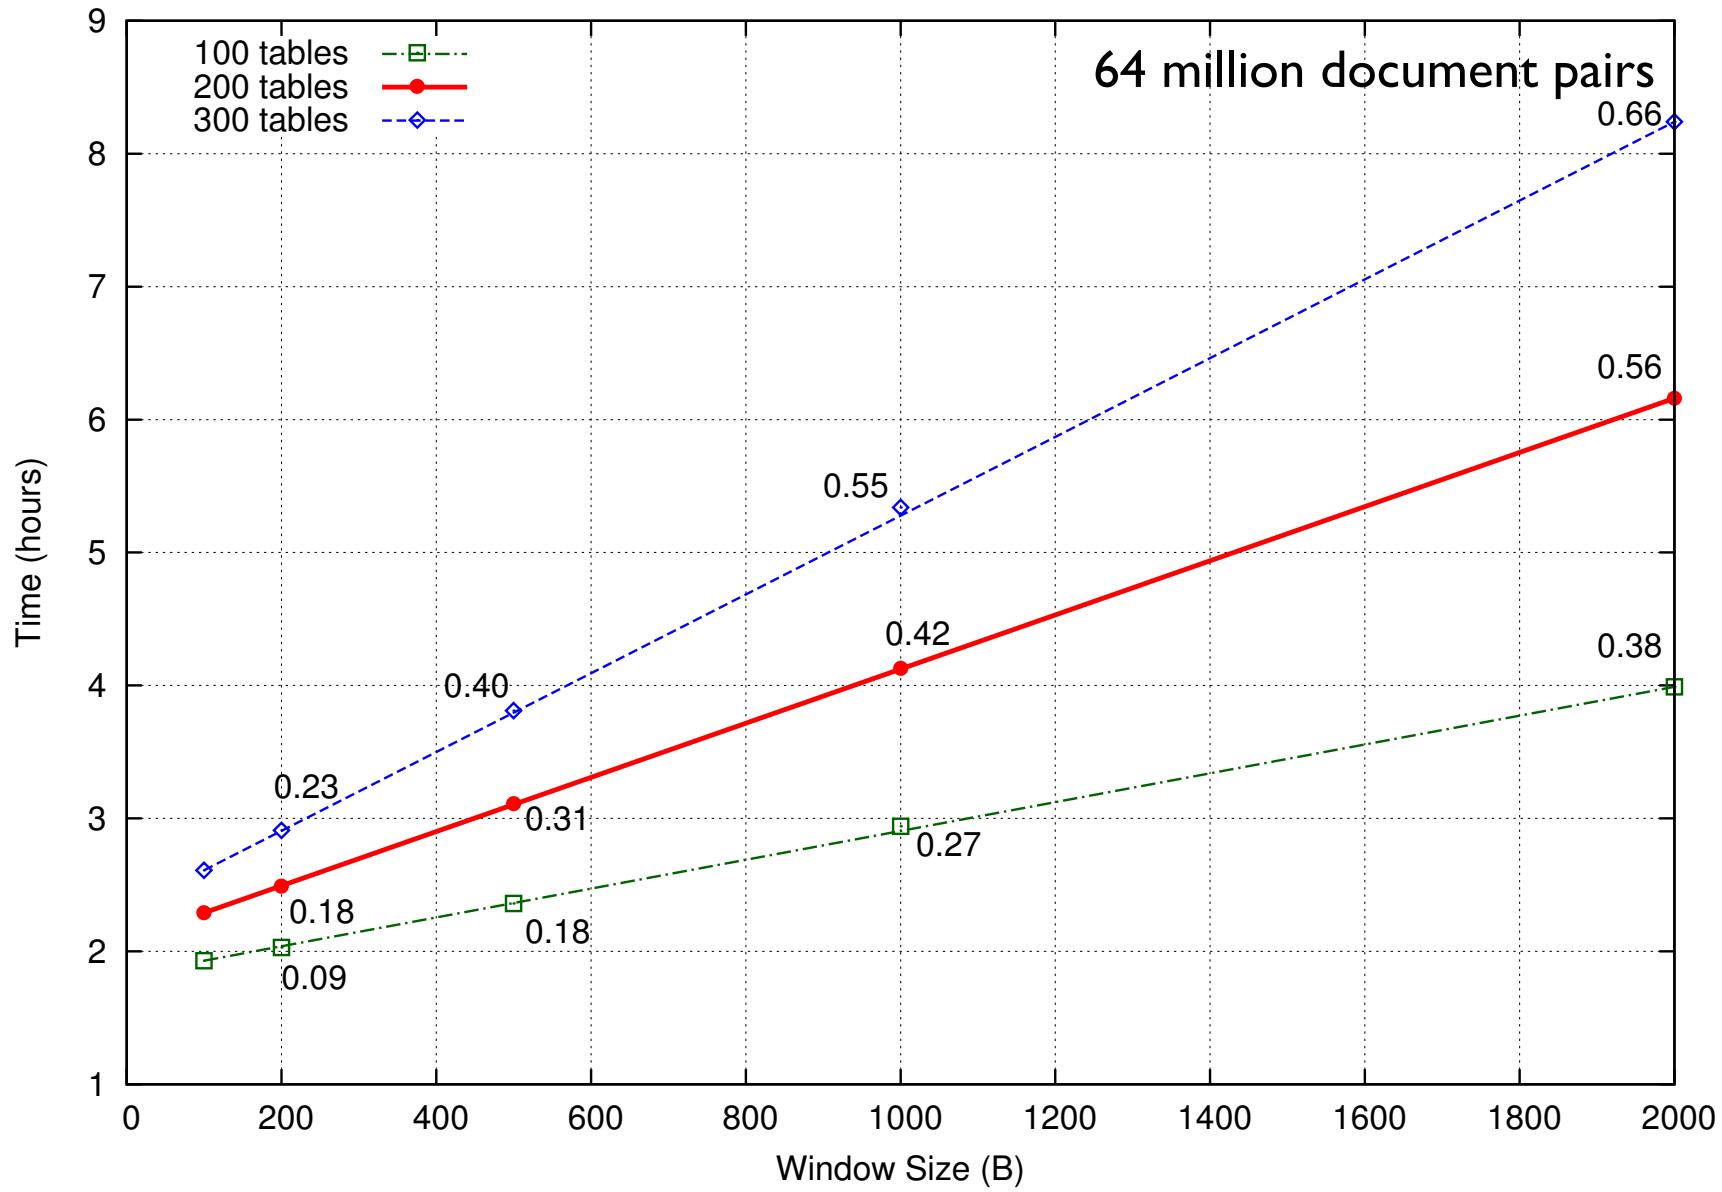
Architecture



Evaluation Setup

- Collection: 3.44m English + 1.47m German Wikipedia
 - Brute force: 5.05 trillion comparisons
- Task: extract similar documents with cosine similarity > 0.3
- Representation: 1000 bit random projections
- Ground truth:
 - Sample ~1000 German Wikipedia articles
- Evaluation metrics:
 - Time
 - Relative cost
 - Recall

Results (16 node cluster)

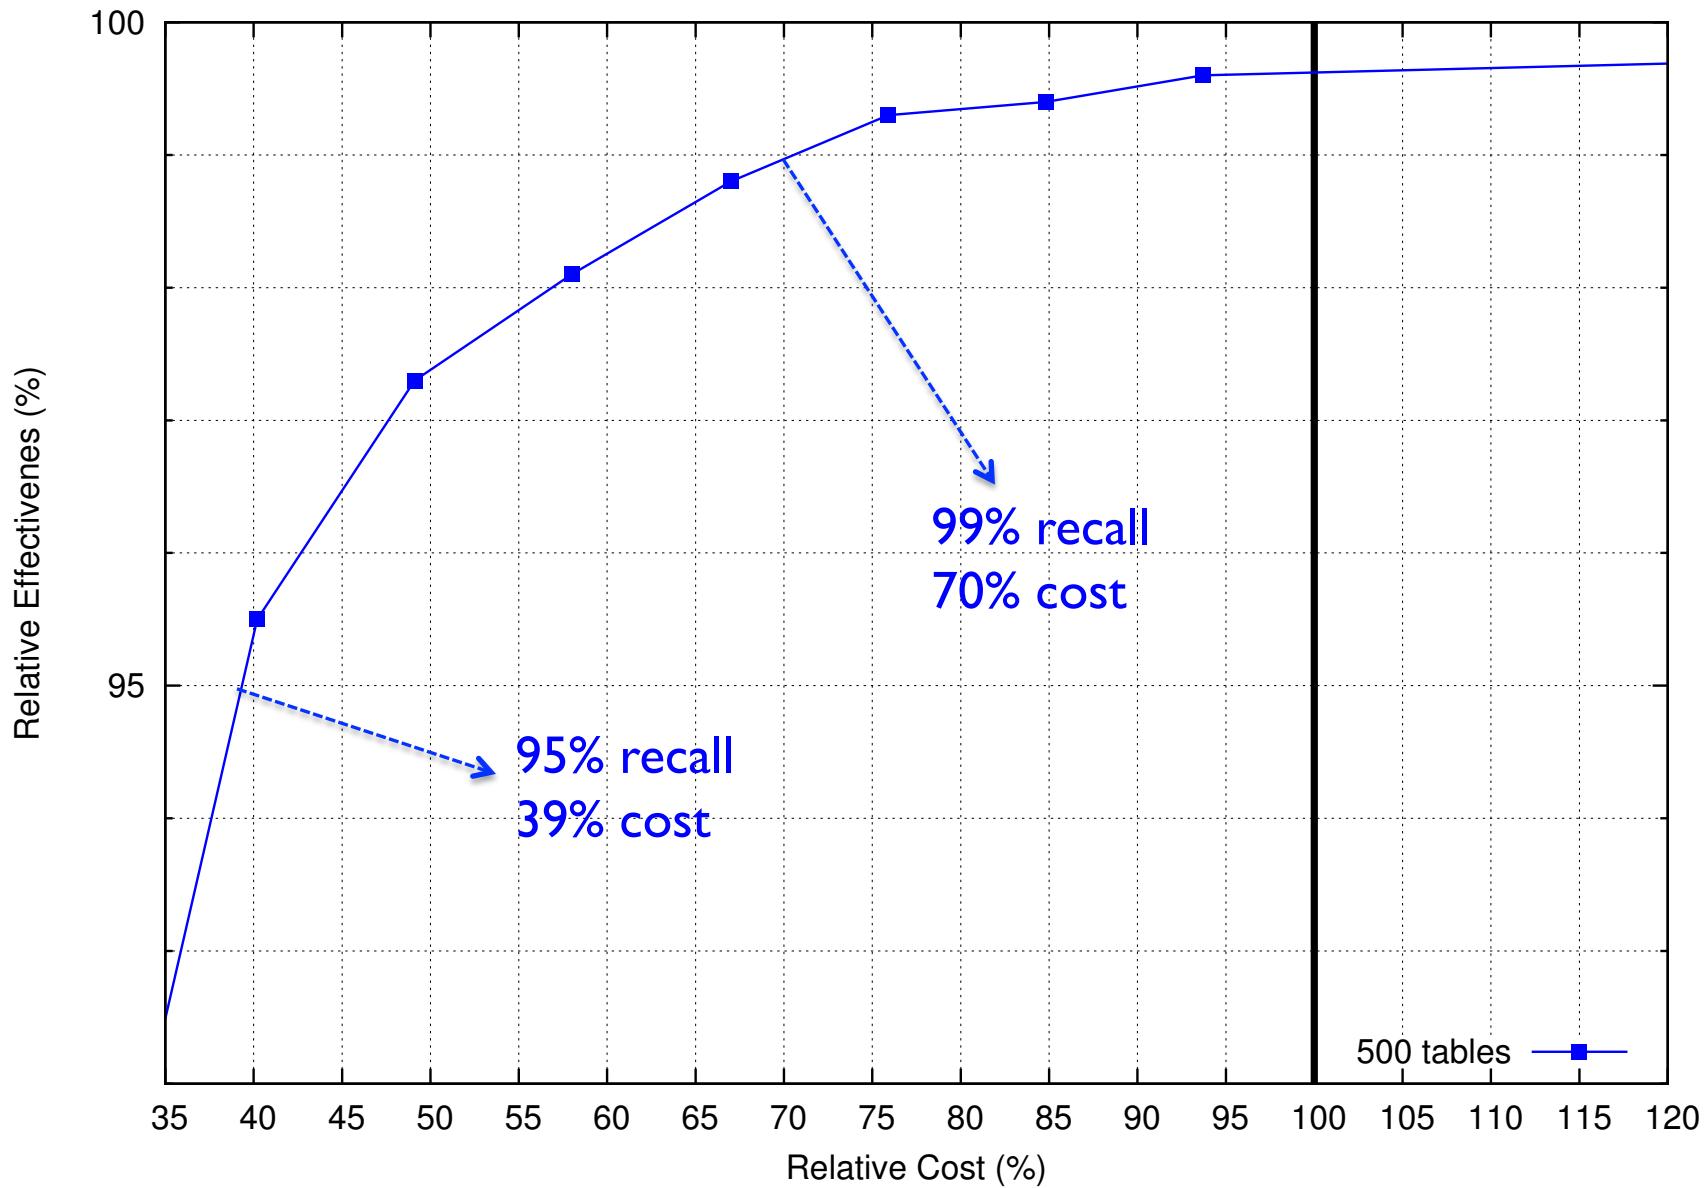


Sources of Error

- Hamming distance computation introduces error!
 - 1000-bit random projections yields ~0.04 absolute error
 - Maximum obtainable recall is 0.76
- Why not just compute cosine similarity?
 - 20 times slower than computing hamming distance

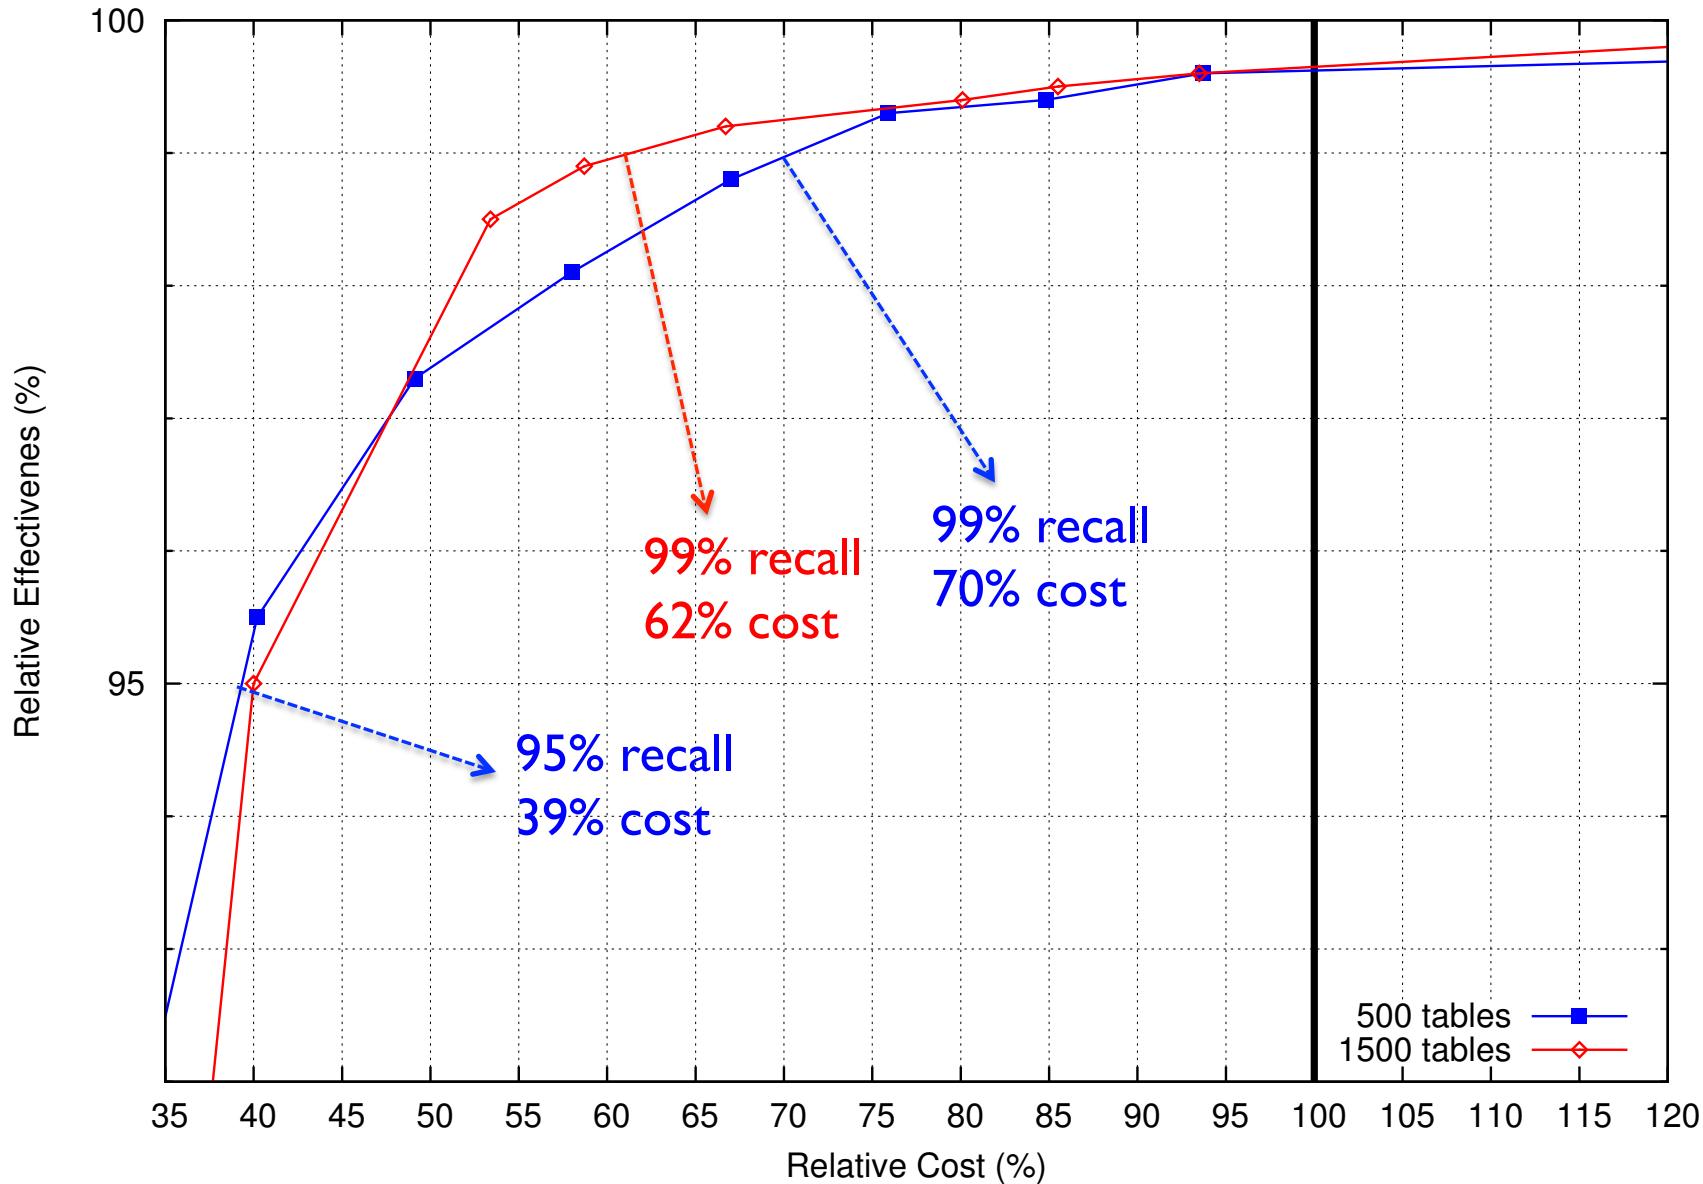
Let's renormalize with respect to hamming distance upper bound...

No Free Lunch!



No Free Lunch!

100% recall? Don't use LSH!



Parallel Sentence Extraction

- Model parallel sentence extraction as a classification problem
 - Given a candidate pair, classify as “parallel” or “not parallel”
- What’s the problem?
 - 64m document pairs yields 400b raw sentence pairs

What's that quote about premature optimization?

Two-Step Classification

- Approach:

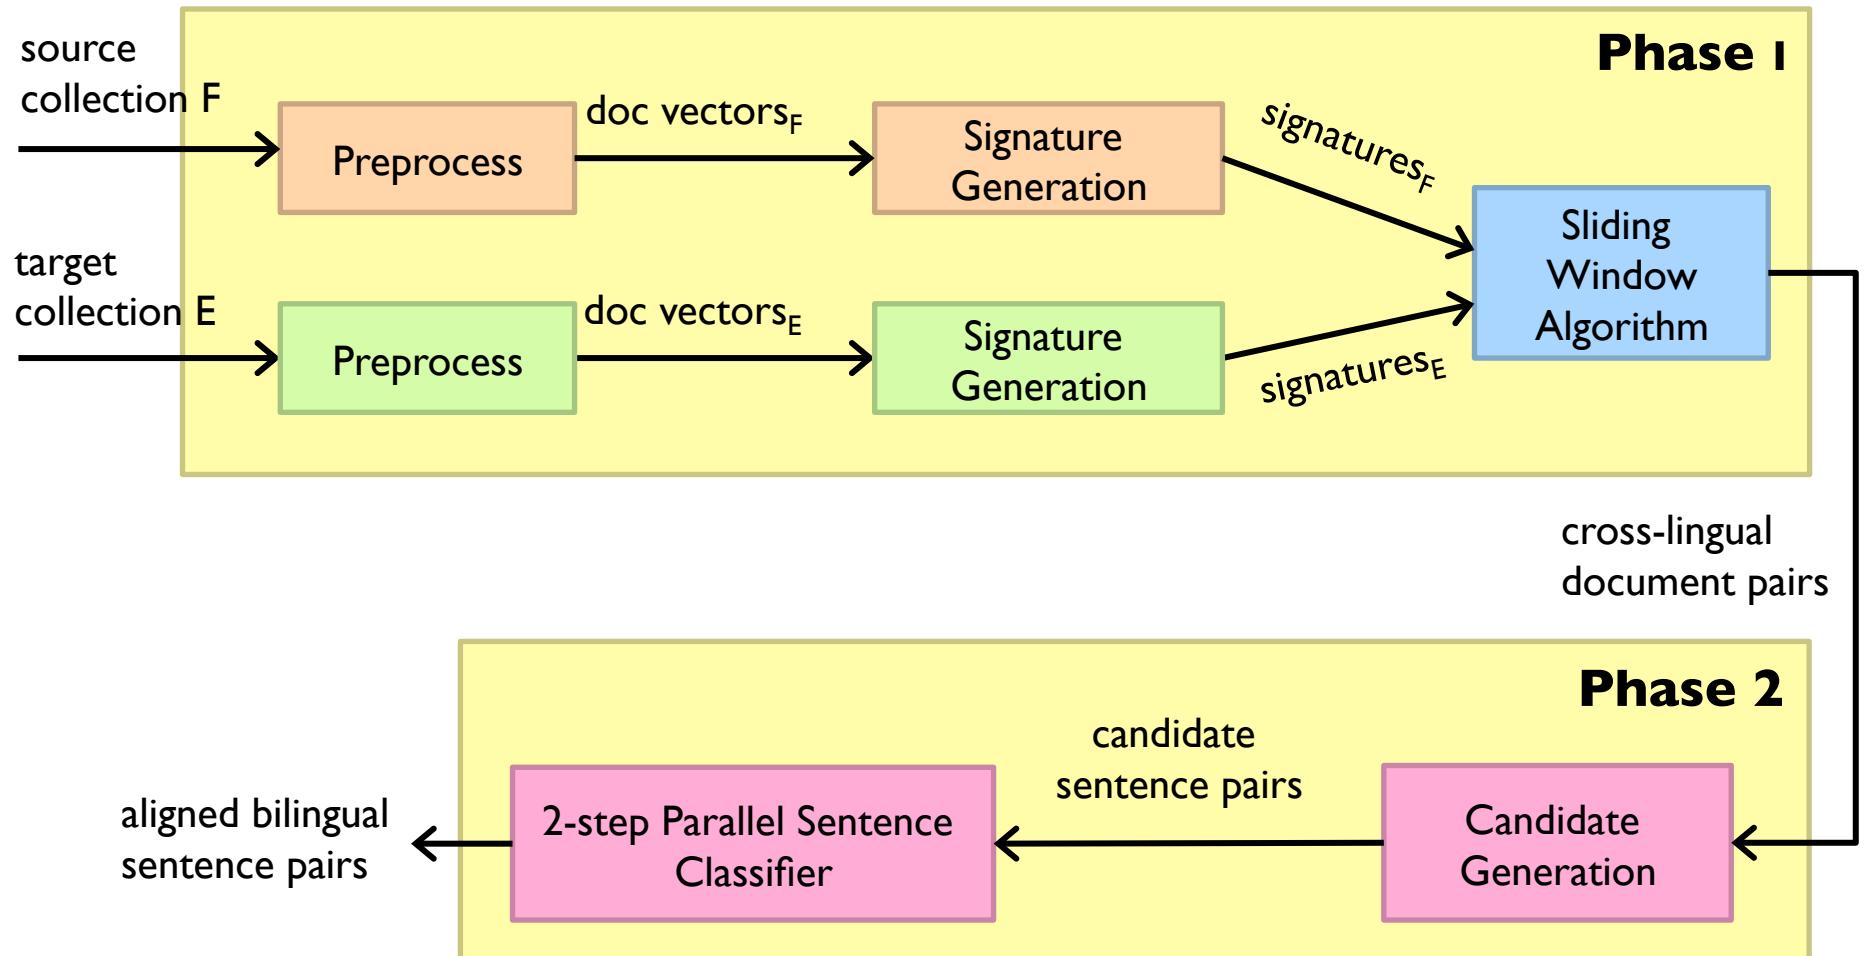
- Apply a few simple heuristics
- “Simple” classifier: efficiently filters out irrelevant pairs
- “Complex” classifier: effectively classifies remaining pairs

Bitext Classifier	Measure	In-domain (Europarl)	Out-of-domain (Wikipedia)	Efficiency (per instance, in μs)
Simple	Recall @ P95	0.81	0.57	27
	Recall @ P80	0.90	0.76	
	Best F-score	0.88	0.78	
Complex	Recall @ P95	0.87	0.55	105
	Recall @ P80	0.94	0.80	
	Best F-score	0.91	0.81	

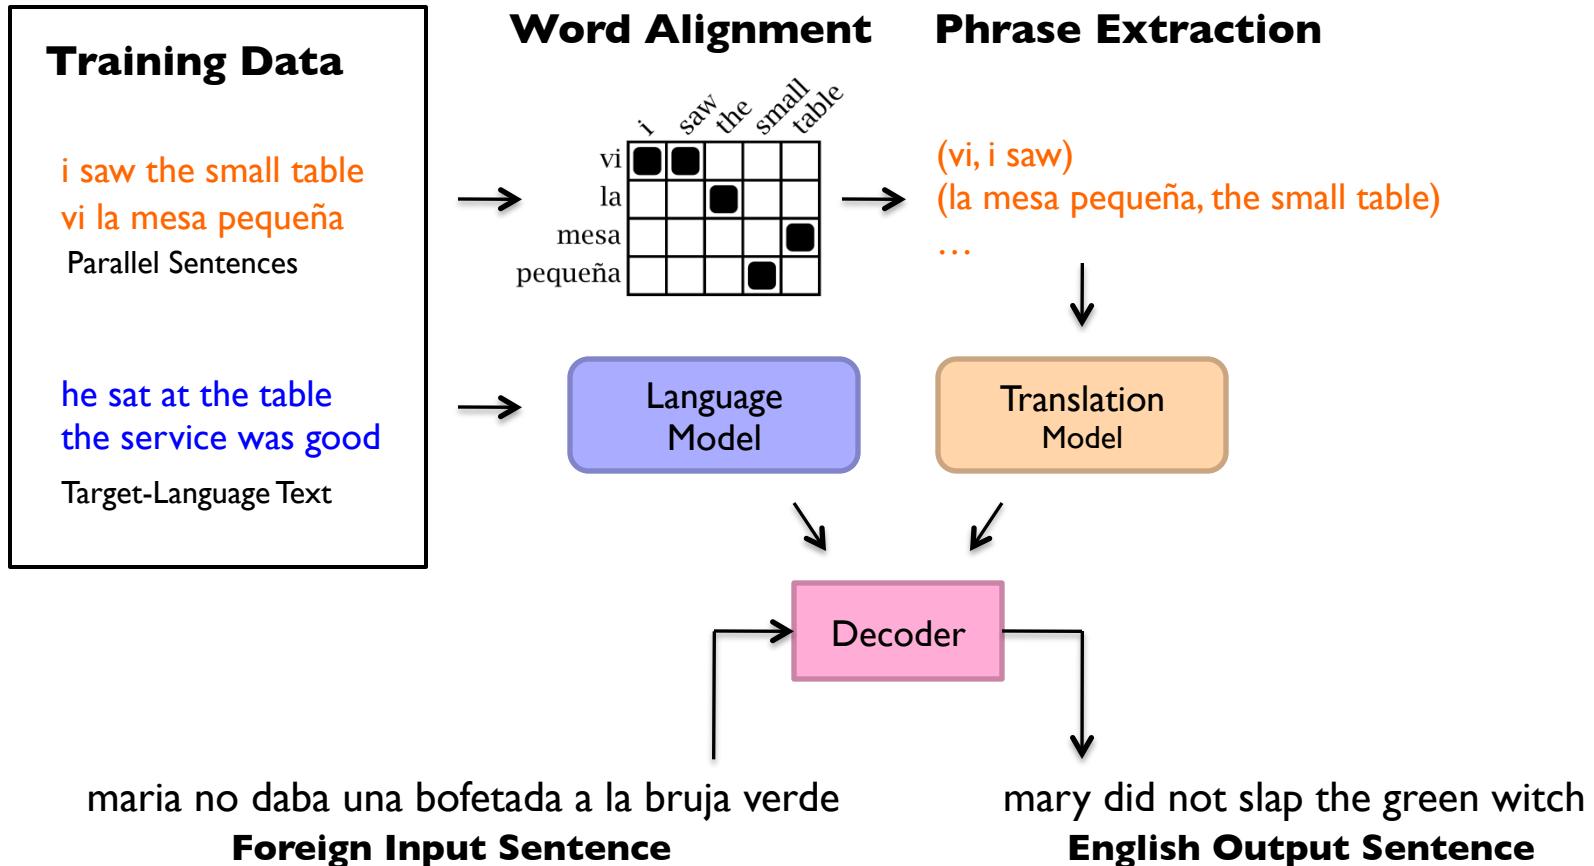
Join

- Algorithm 1
 - Mapper processes both English and foreign collection
 - Load all (d_e, d_f) similarity pairs in memory as side data
 - When encountering relevant document, emit (d_e, d_f) pair as key, list of sentence vectors as value
 - Reducer:
 - Gather sentence vectors, compute Cartesian product and apply classifier
- Algorithm 2
 - Same idea as first, except with stripes pattern
- Which is faster?

Complete Pipeline

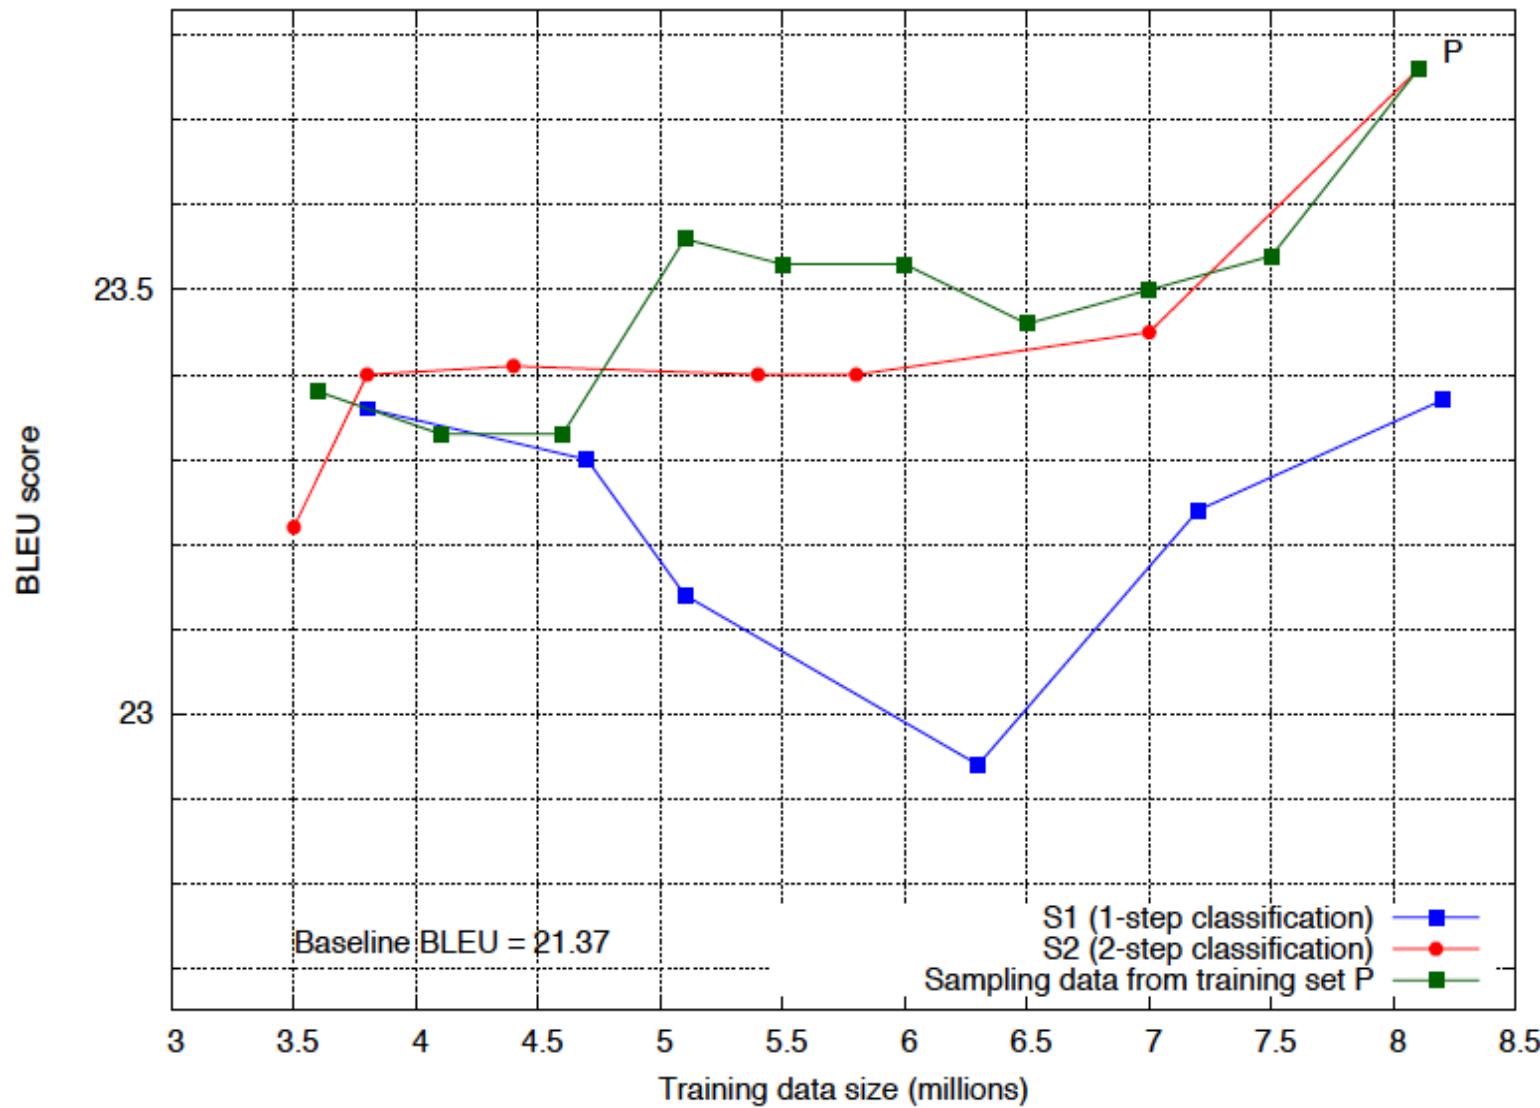


Statistical Machine Translation



$$\hat{e}_1^I = \arg \max_{e_1^I} [P(e_1^I | f_1^J)] = \arg \max_{e_1^I} [P(e_1^I) P(f_1^J | e_1^I)]$$

Evaluation



Today's Agenda

- The problem of similar item detection
- Distances and representations
- Minhash
- Random projections
- End-to-end Wikipedia application

A photograph of a traditional Japanese rock garden. In the foreground, a gravel path is raked into fine, parallel lines. Several large, dark, irregular stones are scattered across the garden. A small, shallow pond is visible in the middle ground, surrounded by more stones and low-lying green plants. In the background, there are more trees and shrubs, and the wooden buildings of a residence are visible behind the garden wall.

Questions?